

Census-Income

February 20, 2019

1 Census-Income (KDD) Data Set

1.1 Data Set Information

This data set contains weighted census data extracted from the 1994 and 1995 Current Population Surveys conducted by the U.S. Census Bureau. The data contains 41 demographic and employment related variables.

One instance per line with comma delimited fields. There are 199523 instances in the data file and 99762 in the test file.

The data was split into train/test in approximately 2/3, 1/3 proportions using MineSet's MIn-dUtil mineset-to-mlc.

Source: <https://archive.ics.uci.edu/ml/datasets/Census-Income+%28KDD%29>

1.2 Attribute Information

More information detailing the meaning of the attributes can be found in the Census Bureau's documentation To make use of the data descriptions at this site, the following mappings to the Census Bureau's internal database column names will be needed:

age 'Age'
class of worker 'ClassOfWorkers'
industry code 'IndustryCode'
occupation code 'OccupationCode'
education 'Education'
wage per hour 'WagePerHour'
enrolled in edu inst last wk 'EnrolledEducation'
marital status 'MaritalStatus'
major industry code 'MajorIndustryCode'
major occupation code 'MajorOccupationCode'
Race 'Race'
hispanic Origin 'HispanicOrigin'
sex 'Sex'
member of a labor union 'LabourUnion'
reason for unemployment 'ReasonUnemployed'
full or part time employment stat 'FullOrPartTime'
capital gains 'CapitalGains'
capital losses 'CapitalLosses'
divdends from stocks 'StockDividends' federal income tax liability
tax filer status 'TaxFilerStat'

region of previous residence 'PrevResidenceRegion'
state of previous residence 'PrevResidenceState'
detailed household and family stat 'HouseholdFamilyStatus'
detailed household summary in household 'HouseholdSummary'
instance weight 'InstanceWeight'
migration code-change in msa 'MigrationCodeChangeMSA'
migration code-change in reg 'MigrationCodeChangeReg'
migration code-move within reg 'MigrationCodeMoveWithinRegion'
live in this house 1 year ago 'LiveInHouse1Y'
migration prev res in sunbelt 'MigPrevResidenceSunbelt'
num persons worked for employer 'NumPersonsWorkedEmployer'
family members under 18 'FamilyMembersU18'
country of birth father 'CountryBirthFather'
country of birth mother 'CountryBirthMother'
country of birth self 'CountryBirthSelf'
citizenship 'Citizenship'
own business or self employed 'OwnBusiness'
fill inc questionnaire for veteran's admin 'QuestionnaireVeteran'
veterans benefits 'VeteranBenefits'
weeks worked in year 'WeeksWorkedInY'

Incomes 'Income'

Note that Incomes have been binned at the \$50K level to present a binary classification problem, much like the original UCI/ ADULT database. The goal field of this data, however, was drawn from the "total person income" field rather than the "adjusted gross income" and may, therefore, behave differently than the original ADULT goal field.

1.3 Problem Statement

The goal of this analysis is to try to find which factors can be used to predict an individual's annual income (higher or lower than 50k USD), and then predict the income level based on these factors.

1.4 Summary

As part of the analysis, I will be going through the following steps:

1- Data Extraction:

Data is available in CSV files that can be downloaded at the source mentioned above.

2- Data Cleaning

3- Exploratory Data Analysis

4- Modeling:

a) Feature Selection: I will use Random Forest Classification for feature selection

b) Model Selection: I will use the selected features and apply Decision Trees and Logistic Regression. I will then proceed with the one with better rates.

5- Oversampling: Given that the data is skewed (only 8% are positive), I will also try oversampling using SMOTE to enhance the recall rates. I will then compare both results (with and without oversampling).

```
In [437]: import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

2 Read the data

In [438]: *#Training and Test data is already split into two CSV's. Hence we will read them sep*

```
columns = ['Age', 'ClassOfWork', 'IndustryCode', 'OccupationCode', 'Education', 'WagePerHour',
           'MaritalStatus', 'MajorIndustryCode', 'MajorOccupationCode', 'Race', 'Hispanic',
           'ReasonUnemployed', 'FullOrPartTime', 'CapitalGains', 'CapitalLosses', 'StockOptions',
           'PrevResidenceRegion', 'PrevResidenceState', 'HouseholdFamilyStatus', 'HouseholdSize',
           'MigrationCodeChangeMSA', 'MigrationCodeChangeReg', 'MigrationCodeMoveWithinMSA',
           'MigPrevResidenceSunbelt', 'NumPersonsWorkedEmployer', 'FamilyMembersU18',
           'CountryBirthMother', 'CountryBirthSelf', 'Citizenship', 'OwnBusiness', 'QCEC',
           'VeteranBenefits', 'WeeksWorkedInY', 'Year', 'Income']
```

```
df = pd.read_csv('Data/census-income.data', header=None)
df.columns = columns
df.drop(['InstanceWeight'], axis=1, inplace=True)
```

```
df_test = pd.read_csv('Data/census-income.test', header=None)
df_test.columns = columns
df_test.drop(['InstanceWeight'], axis=1, inplace=True)
```

In [439]: df.shape

Out[439]: (199523, 41)

In [440]: df_test.shape

Out[440]: (99762, 41)

In [441]: df.head()

```
Out[441]:
```

	Age	ClassOfWork	IndustryCode	OccupationCode	\
0	73	Not in universe	0	0	
1	58	Self-employed-not incorporated	4	34	
2	18	Not in universe	0	0	
3	9	Not in universe	0	0	
4	10	Not in universe	0	0	

	Education	WagePerHour	EnrolledEducation	MaritalStatus	\
0	High school graduate	0	Not in universe	Widowed	
1	Some college but no degree	0	Not in universe	Divorced	
2	10th grade	0	High school	Never married	
3	Children	0	Not in universe	Never married	
4	Children	0	Not in universe	Never married	

	MajorIndustryCode		MajorOccupationCode	\
0	Not in universe or children		Not in universe	
1	Construction	Precision production craft & repair		
2	Not in universe or children		Not in universe	
3	Not in universe or children		Not in universe	
4	Not in universe or children		Not in universe	

	CountryBirthFather	CountryBirthMother	CountryBirthSelf	\
0	United-States	United-States	United-States	
1	United-States	United-States	United-States	
2	Vietnam	Vietnam	Vietnam	
3	United-States	United-States	United-States	
4	United-States	United-States	United-States	

	Citizenship	OwnBusiness	QuestionnaireVeteran	\
0	Native- Born in the United States	0	Not in universe	
1	Native- Born in the United States	0	Not in universe	
2	Foreign born- Not a citizen of U S	0	Not in universe	
3	Native- Born in the United States	0	Not in universe	
4	Native- Born in the United States	0	Not in universe	

	VeteranBenefits	WeeksWorkedInY	Year	Income
0	2	0	95	- 50000.
1	2	52	94	- 50000.
2	2	0	95	- 50000.
3	0	0	94	- 50000.
4	0	0	94	- 50000.

[5 rows x 41 columns]

```
In [442]: #Trim the strings in the data
for i in df.columns:
    if type(df[i][0]) == str:
        df[i] = df[i].apply(lambda x: str(x).strip())
        df_test[i] = df_test[i].apply(lambda x: str(x).strip())

In [443]: #Drop duplicates
df.drop_duplicates(inplace=True)
df_test.drop_duplicates(inplace=True)
```

3 Data Cleaning

```
In [444]: #Check missing data
for i in df.columns:
    print(df[i].unique())
```

```
[73 58 18  9 10 48 42 28 47 34  8 32 51 46 26 13 39 16 35 12 27 56 55  2  1
 37  4 63 25 81 11 30  7 66 84 52  5 36 72 61 41 90 49  6  0 33 57 50 24 17
```

53 40 54 22 29 85 38 76 21 31 74 19 15 3 43 68 71 45 62 23 69 75 44 59 60
 64 65 70 67 78 20 14 83 86 89 77 79 82 80 87 88]
 ['Not in universe' 'Self-employed-not incorporated' 'Private'
 'Local government' 'Federal government' 'Self-employed-incorporated'
 'State government' 'Never worked' 'Without pay']
 [0 4 40 34 43 37 24 39 12 35 45 3 19 29 32 48 33 23 44 36 31 30 41 5 11
 9 42 6 18 50 2 1 26 47 16 14 22 17 7 8 25 46 27 15 13 49 38 21 28 20
 51 10]
 [0 34 10 3 40 26 37 31 12 36 41 22 2 35 25 23 42 8 19 29 27 16 33 13 18
 9 17 39 32 11 30 38 20 7 21 44 24 43 28 4 1 6 45 14 5 15 46]
 ['High school graduate' 'Some college but no degree' '10th grade'
 'Children' 'Bachelors degree(BA AB BS)'
 'Masters degree(MA MS MEng MEd MSW MBA)' 'Less than 1st grade'
 'Associates degree-academic program' '7th and 8th grade'
 '12th grade no diploma' 'Associates degree-occup /vocational'
 'Prof school degree (MD DDS DVM LLB JD)' '5th or 6th grade' '11th grade'
 'Doctorate degree(PhD EdD)' '9th grade' '1st 2nd 3rd or 4th grade']
 [0 1200 876 ..., 3156 2188 1092]
 ['Not in universe' 'High school' 'College or university']
 ['Widowed' 'Divorced' 'Never married' 'Married-civilian spouse present'
 'Separated' 'Married-spouse absent' 'Married-A F spouse present']
 ['Not in universe or children' 'Construction' 'Entertainment'
 'Finance insurance and real estate' 'Education'
 'Business and repair services' 'Manufacturing-nondurable goods'
 'Personal services except private HH' 'Manufacturing-durable goods'
 'Other professional services' 'Mining' 'Transportation' 'Wholesale trade'
 'Public administration' 'Retail trade' 'Social services'
 'Private household services' 'Utilities and sanitary services'
 'Communications' 'Hospital services' 'Medical except hospital'
 'Agriculture' 'Forestry and fisheries' 'Armed Forces']
 ['Not in universe' 'Precision production craft & repair'
 'Professional specialty' 'Executive admin and managerial'
 'Handlers equip cleaners etc' 'Adm support including clerical'
 'Machine operators assmblrs & inspctrs' 'Other service' 'Sales'
 'Private household services' 'Technicians and related support'
 'Transportation and material moving' 'Farming forestry and fishing'
 'Protective services' 'Armed Forces']
 ['White' 'Asian or Pacific Islander' 'Amer Indian Aleut or Eskimo' 'Black'
 'Other']
 ['All other' 'Do not know' 'Central or South American' 'Mexican (Mexicano)'
 'Mexican-American' 'Other Spanish' 'Puerto Rican' 'Cuban' 'Chicano' 'NA']
 ['Female' 'Male']
 ['Not in universe' 'No' 'Yes']
 ['Not in universe' 'Job loser - on layoff' 'Other job loser' 'New entrant'
 'Re-entrant' 'Job leaver']
 ['Not in labor force' 'Children or Armed Forces' 'Full-time schedules'
 'Unemployed full-time' 'Unemployed part- time'
 'PT for non-econ reasons usually FT' 'PT for econ reasons usually PT'

```

'PT for econ reasons usually FT']
[ 0 5178 991 2829 3464 5556 7298 15024 1831 3137 10605 20051
 2538 3908 2407 2050 3103 1086 7688 5013 4386 2414 99999 13550
 2174 4650 4064 914 2354 4787 2009 2597 1055 6097 2635 2105
 3325 6767 2228 2062 3942 27828 9562 2176 7262 2202 2290 1173
 8614 2329 2653 7430 3456 2580 10520 2907 3471 2885 9386 2993
 7896 14084 3818 1409 594 7978 1797 2964 4934 1848 4101 3418
 3432 2774 1424 6849 4687 6418 4508 3674 3411 2936 4416 2346
10566 7443 5455 1151 25236 2463 1455 3781 14344 4865 11678 1471
 5060 114 4931 1506 401 25124 15020 2036 3273 6514 1111 2977
41310 18481 6497 6723 15831 2098 1264 34095 22040 3887 2961 5721
 1090 6360 3800 2387 1731 6612 9472 4594 2601 1140 2227 8530]
[ 0 1590 1977 1669 1719 2444 1421 1848 2205 2149 2001 1902 2090 1573 2415
2377 1876 1602 1740 1974 2339 1887 1258 2597 2603 1408 1980 1721 1816 1340
2788 2174 2042 1485 2489 2129 2457 2051 1762 2057 1672 2258 213 1651 2206
3770 1628 1564 1668 1735 1579 625 4608 2559 2246 4356 1844 2002 2267 3175
1380 2392 1092 1504 2238 2704 2467 810 1539 2824 1741 1870 1944 1825 419
2547 1510 880 1617 1411 1648 323 2282 2352 3004 1755 1429 653 2163 2179
1436 2722 3500 1640 974 1021 2754 1726 3900 2027 772 2231 1138 1594 2465
2519 1956 1911 2472 2201 2080 3683 155]
[ 0 6000 100 ..., 169 1055 7958]
['Nonfiler' 'Head of household' 'Joint both under 65' 'Single'
'Joint both 65+' 'Joint one under 65 & one 65+']
['Not in universe' 'South' 'Northeast' 'Midwest' 'West' 'Abroad']
['Not in universe' 'Arkansas' 'Utah' 'Michigan' 'Minnesota' 'Alaska'
'Kansas' 'Indiana' '?' 'Massachusetts' 'New Mexico' 'Nevada' 'Tennessee'
'Colorado' 'Abroad' 'Kentucky' 'California' 'Arizona' 'North Carolina'
'Connecticut' 'Florida' 'Vermont' 'Maryland' 'Oklahoma' 'Oregon' 'Ohio'
'South Carolina' 'Texas' 'Montana' 'Wyoming' 'Georgia' 'Pennsylvania'
'Iowa' 'New Hampshire' 'Missouri' 'Alabama' 'North Dakota' 'New Jersey'
'Louisiana' 'West Virginia' 'Delaware' 'Illinois' 'Maine' 'Wisconsin'
'New York' 'Idaho' 'District of Columbia' 'South Dakota' 'Nebraska'
'Virginia' 'Mississippi']
['Other Rel 18+ ever marr not in subfamily' 'Householder'
'Child 18+ never marr Not in a subfamily'
'Child <18 never marr not in subfamily' 'Spouse of householder'
'Secondary individual' 'Other Rel 18+ never marr not in subfamily'
'Nonfamily householder' 'Grandchild <18 never marr not in subfamily'
'Grandchild <18 never marr child of subfamily RP'
'Child 18+ ever marr Not in a subfamily'
'Child 18+ never marr RP of subfamily' 'Child 18+ spouse of subfamily RP'
'Other Rel <18 never marr child of subfamily RP'
'Child under 18 of RP of unrel subfamily'
'Grandchild 18+ never marr not in subfamily'
'Child 18+ ever marr RP of subfamily'
'Other Rel 18+ ever marr RP of subfamily' 'RP of unrelated subfamily'
'Other Rel 18+ spouse of subfamily RP'
'Other Rel <18 never marr not in subfamily'

```

'Other Rel <18 spouse of subfamily RP' 'In group quarters'
 'Grandchild 18+ spouse of subfamily RP'
 'Other Rel 18+ never marr RP of subfamily'
 'Child <18 never marr RP of subfamily'
 'Child <18 ever marr not in subfamily'
 'Other Rel <18 ever marr RP of subfamily'
 'Grandchild 18+ ever marr not in subfamily'
 'Child <18 spouse of subfamily RP' 'Spouse of RP of unrelated subfamily'
 'Other Rel <18 never married RP of subfamily'
 'Grandchild 18+ never marr RP of subfamily'
 'Grandchild 18+ ever marr RP of subfamily'
 'Child <18 ever marr RP of subfamily'
 'Other Rel <18 ever marr not in subfamily'
 'Grandchild <18 never marr RP of subfamily'
 'Grandchild <18 ever marr not in subfamily']
 ['Other relative of householder' 'Householder' 'Child 18 or older'
 'Child under 18 never married' 'Spouse of householder'
 'Nonrelative of householder' 'Group Quarters- Secondary individual'
 'Child under 18 ever married']
 ['?' 'MSA to MSA' 'Nonmover' 'NonMSA to nonMSA' 'Not in universe'
 'Not identifiable' 'Abroad to MSA' 'MSA to nonMSA' 'Abroad to nonMSA'
 'NonMSA to MSA']
 ['?' 'Same county' 'Nonmover' 'Different region'
 'Different county same state' 'Not in universe'
 'Different division same region' 'Abroad' 'Different state same division']
 ['?' 'Same county' 'Nonmover' 'Different state in South'
 'Different county same state' 'Not in universe'
 'Different state in Northeast' 'Abroad' 'Different state in Midwest'
 'Different state in West']
 ['Not in universe under 1 year old' 'No' 'Yes']
 ['?' 'Yes' 'Not in universe' 'No']
 [0 1 6 4 5 3 2]
 ['Not in universe' 'Both parents present' 'Mother only present'
 'Neither parent present' 'Father only present']
 ['United-States' 'Vietnam' 'Philippines' '?' 'Columbia' 'Germany' 'Mexico'
 'Japan' 'Peru' 'Dominican-Republic' 'South Korea' 'Cuba' 'El-Salvador'
 'Canada' 'Scotland' 'Outlying-U S (Guam USVI etc)' 'Italy' 'Guatemala'
 'Ecuador' 'Puerto-Rico' 'Cambodia' 'China' 'Poland' 'Nicaragua' 'Taiwan'
 'England' 'Ireland' 'Hungary' 'Yugoslavia' 'Trinidad&Tobago' 'Jamaica'
 'Honduras' 'Portugal' 'Iran' 'France' 'India' 'Hong Kong' 'Haiti' 'Greece'
 'Holand-Netherlands' 'Thailand' 'Laos' 'Panama']
 ['United-States' 'Vietnam' '?' 'Columbia' 'Mexico' 'El-Salvador' 'Peru'
 'Puerto-Rico' 'Cuba' 'Philippines' 'Dominican-Republic' 'Germany'
 'England' 'Guatemala' 'Scotland' 'Portugal' 'Italy' 'Ecuador' 'Yugoslavia'
 'China' 'Poland' 'Hungary' 'Nicaragua' 'Taiwan' 'Ireland' 'Canada'
 'South Korea' 'Trinidad&Tobago' 'Jamaica' 'Honduras' 'Iran' 'France'
 'Cambodia' 'India' 'Hong Kong' 'Haiti' 'Japan' 'Greece'
 'Holand-Netherlands' 'Thailand' 'Panama' 'Laos'

```

'Outlying-U S (Guam USVI etc)']
['United-States' 'Vietnam' '?' 'Columbia' 'Mexico' 'Peru' 'Cuba'
'Philippines' 'Dominican-Republic' 'El-Salvador' 'Canada' 'Scotland'
'Portugal' 'Guatemala' 'Ecuador' 'Germany' 'Outlying-U S (Guam USVI etc)'
'Puerto-Rico' 'Italy' 'China' 'Poland' 'Nicaragua' 'Taiwan' 'England'
'Ireland' 'South Korea' 'Trinidad&Tobago' 'Jamaica' 'Honduras' 'Iran'
'Hungary' 'France' 'Cambodia' 'India' 'Hong Kong' 'Japan' 'Haiti'
'Holand-Netherlands' 'Greece' 'Thailand' 'Panama' 'Yugoslavia' 'Laos']
['Native- Born in the United States' 'Foreign born- Not a citizen of U S'
'Foreign born- U S citizen by naturalization'
'Native- Born abroad of American Parent(s)'
'Native- Born in Puerto Rico or U S Outlying']
[0 2 1]
['Not in universe' 'No' 'Yes']
[2 0 1]
[ 0 52 30 49 32 15 38 48  9 24 50 10 45 43  4 26 40 20  6 12 51  1  8 39 13
16 34 14 36 44 22 41 46 28 23 35 25 17 11 37  5 42 29  2 21 19 47  3 27  7
18 33 31]
[95 94]
['- 50000.' '50000+.']

```

```

In [445]: #Some missing data is represented as '?', others are 'Not in universe'
df.replace("?", np.nan, inplace=True)
df.replace("Not in universe", np.nan, inplace=True)
df_test.replace("?", np.nan, inplace=True)
df_test.replace("Not in universe", np.nan, inplace=True)

```

```

In [446]: df.isnull().sum()

```

```

Out[446]: Age                                0
ClassOfWorker                             54165
IndustryCode                             0
OccupationCode                           0
Education                                0
WagePerHour                              0
EnrolledEducation                       142243
MaritalStatus                           0
MajorIndustryCode                       0
MajorOccupationCode                     54548
Race                                    0
HispanicOrigin                          0
Sex                                      0
LabourUnion                             133840
ReasonUnemployed                       146884
FullOrPartTime                          0
CapitalGains                            0
CapitalLosses                           0

```


StockDividends	0
TaxFilerStat	0
PrevResidenceRegion	137492
PrevResidenceState	138190
HouseholdFamilyStatus	0
HouseholdSummary	0
MigrationCodeChangeMSA	75288
MigrationCodeChangeReg	75288
MigrationCodeMoveWithinRegion	75288
LiveInHouse1Y	0
MigPrevResidenceSunbelt	137492
NumPersonsWorkedEmployer	0
FamilyMembersU18	134959
CountryBirthFather	6383
CountryBirthMother	5810
CountryBirthSelf	3322
Citizenship	0
OwnBusiness	0
QuestionnaireVeteran	150931
VeteranBenefits	0
WeeksWorkedInY	0
Year	0
Income	0
dtype:	int64

In [447]: *#Delete the columns with almost completely missing values*

```
df.drop(['EnrolledEducation', 'LabourUnion', 'ReasonUnemployed', 'PrevResidenceRegion',
        'MigPrevResidenceSunbelt', 'QuestionnaireVeteran', 'FamilyMembersU18', 'MigrationCodeChangeReg', 'MigrationCodeMoveWithinRegion'], axis=1, inplace=True)

df_test.drop(['EnrolledEducation', 'LabourUnion', 'ReasonUnemployed', 'PrevResidenceRegion',
              'MigPrevResidenceSunbelt', 'QuestionnaireVeteran', 'FamilyMembersU18', 'MigrationCodeChangeReg', 'MigrationCodeMoveWithinRegion'], axis=1, inplace=True)
```

In [448]: *#Drop Rows with few missing column values*

```
df.dropna(subset=['CountryBirthFather', 'CountryBirthMother', 'CountryBirthSelf'], inplace=True)
df_test.dropna(subset=['CountryBirthFather', 'CountryBirthMother', 'CountryBirthSelf'], inplace=True)
```

In [449]: df.isnull().sum()

Age	0
ClassOfWork	49901
IndustryCode	0
OccupationCode	0
Education	0
WagePerHour	0
MaritalStatus	0
MajorIndustryCode	0
MajorOccupationCode	50258

```

Race                                0
HispanicOrigin                      0
Sex                                 0
FullOrPartTime                      0
CapitalGains                        0
CapitalLosses                       0
StockDividends                     0
TaxFilerStat                        0
HouseholdFamilyStatus               0
HouseholdSummary                    0
LiveInHouse1Y                       0
NumPersonsWorkedEmployer            0
CountryBirthFather                  0
CountryBirthMother                  0
CountryBirthSelf                    0
Citizenship                         0
OwnBusiness                         0
VeteranBenefits                     0
WeeksWorkedInY                     0
Year                                0
Income                              0
dtype: int64

```

```

In [450]: #ClassOfWork and MajorOccupationCode are missing for those who do not work. I will
df[['ClassOfWork', 'MajorOccupationCode']] = df[['ClassOfWork', 'MajorOccupationCode']]
df_test[['ClassOfWork', 'MajorOccupationCode']] = df_test[['ClassOfWork', 'MajorOccupationCode']]

```

```

In [451]: df['Income>50k'] = np.where(df['Income'] == '- 50000.', 0, 1)
df.drop('Income', axis=1, inplace=True)

df_test['Income>50k'] = np.where(df_test['Income'] == '- 50000.', 0, 1)
df_test.drop('Income', axis=1, inplace=True)

```

```

In [452]: df.head(100)

```

```

Out[452]:
   Age  ClassOfWork  IndustryCode  OccupationCode  \
0    73           NA             0              0
1    58  Self-employed-not incorporated         4          34
2    18           NA             0              0
3     9           NA             0              0
4    10           NA             0              0
5    48        Private          40             10
6    42        Private          34              3
7    28        Private           4            40
8    47  Local government          43            26
9    34        Private           4            37
10   8           NA             0              0
12   51        Private           4            34
13   46        Private          37            31

```

14	26	Private	24	12
15	13	NA	0	0
16	47	Private	39	36
17	39	NA	0	0
18	16	NA	0	0
19	35	Private	12	41
20	12	NA	0	0
21	27	Self-employed-not incorporated	4	34
22	56	Private	35	22
23	46	Private	45	12
24	55	NA	0	0
25	2	NA	0	0
27	37	Private	3	34
28	4	NA	0	0
29	37	Private	4	2
30	63	Private	19	35
31	34	Federal government	29	25
..
74	51	Private	32	18
75	10	NA	0	0
76	48	State government	43	9
77	27	Private	33	16
78	24	NA	0	0
79	17	NA	0	0
80	58	Self-employed-not incorporated	35	17
81	2	NA	0	0
82	53	Private	4	34
83	33	Private	45	23
84	7	NA	0	0
85	40	Private	19	42
86	25	Private	33	29
88	54	NA	0	0
89	22	Private	34	26
90	29	NA	0	0
91	9	NA	0	0
93	40	Private	29	17
94	38	NA	0	0
95	7	NA	0	0
96	49	State government	29	26
98	7	NA	0	0
99	21	Private	33	3
100	52	NA	0	0
101	34	Local government	43	10
102	6	NA	0	0
103	31	Private	5	39
104	74	NA	0	0
105	0	NA	0	0
106	19	Private	39	32

	Education	WagePerHour \
0	High school graduate	0
1	Some college but no degree	0
2	10th grade	0
3	Children	0
4	Children	0
5	Some college but no degree	1200
6	Bachelors degree(BA AB BS)	0
7	High school graduate	0
8	Some college but no degree	876
9	Some college but no degree	0
10	Children	0
12	Some college but no degree	0
13	High school graduate	0
14	Bachelors degree(BA AB BS)	0
15	Children	0
16	Bachelors degree(BA AB BS)	0
17	10th grade	0
18	10th grade	0
19	High school graduate	0
20	Children	0
21	Some college but no degree	0
22	Some college but no degree	500
23	Masters degree(MA MS MEng MEd MSW MBA)	0
24	Some college but no degree	0
25	Children	0
27	Some college but no degree	0
28	Children	0
29	Bachelors degree(BA AB BS)	0
30	Less than 1st grade	0
31	Some college but no degree	0
..
74	Some college but no degree	0
75	Children	0
76	Some college but no degree	0
77	Some college but no degree	0
78	High school graduate	0
79	7th and 8th grade	0
80	Prof school degree (MD DDS DVM LLB JD)	0
81	Children	0
82	10th grade	0
83	Bachelors degree(BA AB BS)	0
84	Children	0
85	5th or 6th grade	0
86	High school graduate	0
88	High school graduate	0
89	High school graduate	0

90	Some college but no degree	0
91	Children	0
93	High school graduate	0
94	Less than 1st grade	0
95	Children	0
96	11th grade	0
98	Children	0
99	Some college but no degree	0
100	7th and 8th grade	0
101	Bachelors degree(BA AB BS)	0
102	Children	0
103	11th grade	0
104	High school graduate	0
105	Children	0
106	Some college but no degree	0

	MaritalStatus	MajorIndustryCode \
0	Widowed	Not in universe or children
1	Divorced	Construction
2	Never married	Not in universe or children
3	Never married	Not in universe or children
4	Never married	Not in universe or children
5	Married-civilian spouse present	Entertainment
6	Married-civilian spouse present	Finance insurance and real estate
7	Never married	Construction
8	Married-civilian spouse present	Education
9	Married-civilian spouse present	Construction
10	Never married	Not in universe or children
12	Married-civilian spouse present	Construction
13	Divorced	Business and repair services
14	Never married	Manufacturing-nondurable goods
15	Never married	Not in universe or children
16	Never married	Personal services except private HH
17	Married-civilian spouse present	Not in universe or children
18	Never married	Not in universe or children
19	Married-civilian spouse present	Manufacturing-durable goods
20	Never married	Not in universe or children
21	Married-civilian spouse present	Construction
22	Married-civilian spouse present	Finance insurance and real estate
23	Married-civilian spouse present	Other professional services
24	Married-civilian spouse present	Not in universe or children
25	Never married	Not in universe or children
27	Married-civilian spouse present	Mining
28	Never married	Not in universe or children
29	Never married	Construction
30	Married-civilian spouse present	Manufacturing-nondurable goods
31	Married-civilian spouse present	Transportation
..

74	Married-civilian spouse present	Wholesale trade
75	Never married	Not in universe or children
76	Married-civilian spouse present	Education
77	Never married	Retail trade
78	Married-civilian spouse present	Not in universe or children
79	Never married	Not in universe or children
80	Married-civilian spouse present	Finance insurance and real estate
81	Never married	Not in universe or children
82	Married-civilian spouse present	Construction
83	Never married	Other professional services
84	Never married	Not in universe or children
85	Married-civilian spouse present	Manufacturing-nondurable goods
86	Never married	Retail trade
88	Married-civilian spouse present	Not in universe or children
89	Married-civilian spouse present	Finance insurance and real estate
90	Married-civilian spouse present	Not in universe or children
91	Never married	Not in universe or children
93	Married-civilian spouse present	Transportation
94	Never married	Not in universe or children
95	Never married	Not in universe or children
96	Married-civilian spouse present	Transportation
98	Never married	Not in universe or children
99	Never married	Retail trade
100	Married-civilian spouse present	Not in universe or children
101	Married-civilian spouse present	Education
102	Never married	Not in universe or children
103	Married-civilian spouse present	Manufacturing-durable goods
104	Married-civilian spouse present	Not in universe or children
105	Never married	Not in universe or children
106	Never married	Personal services except private HH

	MajorOccupationCode	Race \
0	NA	White
1	Precision production craft & repair	White
2	NA	Asian or Pacific Islander
3	NA	White
4	NA	White
5	Professional specialty	Amer Indian Aleut or Eskimo
6	Executive admin and managerial	White
7	Handlers equip cleaners etc	White
8	Adm support including clerical	White
9	Machine operators assmblrs & inspctrs	White
10	NA	White
12	Precision production craft & repair	White
13	Other service	White
14	Professional specialty	White
15	NA	Black
16	Machine operators assmblrs & inspctrs	White

17		NA	White
18		NA	White
19	Handlers equip cleaners etc		White
20		NA	Other
21	Precision production craft & repair		White
22	Adm support including clerical		White
23	Professional specialty		White
24		NA	Asian or Pacific Islander
25		NA	White
27	Precision production craft & repair		White
28		NA	White
29	Executive admin and managerial		White
30	Precision production craft & repair		Other
31	Adm support including clerical		White
..	
74		Sales	White
75		NA	White
76	Professional specialty		White
77		Sales	White
78		NA	Black
79		NA	White
80		Sales	White
81		NA	Black
82	Precision production craft & repair		White
83	Adm support including clerical		Black
84		NA	White
85	Handlers equip cleaners etc		White
86	Other service		White
88		NA	Asian or Pacific Islander
89	Adm support including clerical		White
90		NA	Black
91		NA	White
93		Sales	White
94		NA	White
95		NA	White
96	Adm support including clerical		White
98		NA	White
99	Executive admin and managerial		White
100		NA	White
101	Professional specialty		White
102		NA	White
103	Transportation and material moving		White
104		NA	White
105		NA	White
106	Other service		White

...	NumPersonsWorkedEmployer	CountryBirthFather	CountryBirthMother	\
0	...	0	United-States	United-States

1	...	1	United-States	United-States
2	...	0	Vietnam	Vietnam
3	...	0	United-States	United-States
4	...	0	United-States	United-States
5	...	1	Philippines	United-States
6	...	6	United-States	United-States
7	...	4	United-States	United-States
8	...	5	United-States	United-States
9	...	6	United-States	United-States
10	...	0	United-States	United-States
12	...	3	United-States	United-States
13	...	6	Columbia	Columbia
14	...	6	United-States	United-States
15	...	0	United-States	United-States
16	...	6	Germany	United-States
17	...	0	Mexico	Mexico
18	...	0	United-States	El-Salvador
19	...	4	United-States	United-States
20	...	0	United-States	United-States
21	...	6	United-States	United-States
22	...	2	United-States	United-States
23	...	1	United-States	United-States
24	...	0	Japan	United-States
25	...	0	United-States	United-States
27	...	3	United-States	United-States
28	...	0	United-States	United-States
29	...	6	United-States	United-States
30	...	6	Mexico	Mexico
31	...	6	United-States	United-States
..
74	...	4	United-States	United-States
75	...	0	United-States	Mexico
76	...	6	United-States	United-States
77	...	6	United-States	United-States
78	...	6	United-States	United-States
79	...	0	United-States	United-States
80	...	1	United-States	United-States
81	...	0	United-States	United-States
82	...	1	United-States	United-States
83	...	6	United-States	United-States
84	...	0	Germany	United-States
85	...	3	Mexico	Mexico
86	...	3	United-States	United-States
88	...	0	Vietnam	Vietnam
89	...	3	United-States	United-States
90	...	6	United-States	United-States
91	...	0	United-States	United-States
93	...	3	United-States	United-States

94	...	0	United-States	United-States
95	...	0	United-States	United-States
96	...	6	United-States	United-States
98	...	0	United-States	United-States
99	...	6	United-States	United-States
100	...	0	United-States	United-States
101	...	6	United-States	United-States
102	...	0	United-States	United-States
103	...	2	United-States	United-States
104	...	0	United-States	United-States
105	...	0	Mexico	Mexico
106	...	3	United-States	United-States

	CountryBirthSelf	Citizenship	OwnBusiness	\
0	United-States	Native- Born in the United States	0	
1	United-States	Native- Born in the United States	0	
2	Vietnam	Foreign born- Not a citizen of U S	0	
3	United-States	Native- Born in the United States	0	
4	United-States	Native- Born in the United States	0	
5	United-States	Native- Born in the United States	2	
6	United-States	Native- Born in the United States	0	
7	United-States	Native- Born in the United States	0	
8	United-States	Native- Born in the United States	0	
9	United-States	Native- Born in the United States	0	
10	United-States	Native- Born in the United States	0	
12	United-States	Native- Born in the United States	0	
13	Columbia	Foreign born- Not a citizen of U S	0	
14	United-States	Native- Born in the United States	0	
15	United-States	Native- Born in the United States	0	
16	United-States	Native- Born in the United States	0	
17	Mexico	Foreign born- Not a citizen of U S	0	
18	United-States	Native- Born in the United States	0	
19	United-States	Native- Born in the United States	0	
20	United-States	Native- Born in the United States	0	
21	United-States	Native- Born in the United States	1	
22	United-States	Native- Born in the United States	2	
23	United-States	Native- Born in the United States	0	
24	United-States	Native- Born in the United States	0	
25	United-States	Native- Born in the United States	0	
27	United-States	Native- Born in the United States	0	
28	United-States	Native- Born in the United States	0	
29	United-States	Native- Born in the United States	0	
30	Mexico	Foreign born- Not a citizen of U S	0	
31	United-States	Native- Born in the United States	0	
..	
74	United-States	Native- Born in the United States	0	
75	United-States	Native- Born in the United States	0	
76	United-States	Native- Born in the United States	2	

77	United-States	Native- Born in the United States	0
78	United-States	Native- Born in the United States	0
79	United-States	Native- Born in the United States	0
80	United-States	Native- Born in the United States	0
81	United-States	Native- Born in the United States	0
82	United-States	Native- Born in the United States	0
83	United-States	Native- Born in the United States	0
84	United-States	Native- Born in the United States	0
85	Mexico	Foreign born- Not a citizen of U S	0
86	United-States	Native- Born in the United States	0
88	Vietnam	Foreign born- Not a citizen of U S	0
89	United-States	Native- Born in the United States	0
90	United-States	Native- Born in the United States	0
91	United-States	Native- Born in the United States	0
93	United-States	Native- Born in the United States	0
94	United-States	Native- Born in the United States	0
95	United-States	Native- Born in the United States	0
96	United-States	Native- Born in the United States	0
98	United-States	Native- Born in the United States	0
99	United-States	Native- Born in the United States	2
100	United-States	Native- Born in the United States	0
101	United-States	Native- Born in the United States	0
102	United-States	Native- Born in the United States	0
103	United-States	Native- Born in the United States	0
104	United-States	Native- Born in the United States	0
105	United-States	Native- Born in the United States	0
106	United-States	Native- Born in the United States	0

	VeteranBenefits	WeeksWorkedInY	Year	Income>50k
0	2	0	95	0
1	2	52	94	0
2	2	0	95	0
3	0	0	94	0
4	0	0	94	0
5	2	52	95	0
6	2	52	94	0
7	2	30	95	0
8	2	52	95	0
9	2	52	94	0
10	0	0	94	0
12	2	52	94	0
13	2	52	94	0
14	2	52	95	0
15	0	0	94	0
16	2	52	95	0
17	2	0	94	0
18	2	0	95	0
19	2	49	95	0

20	0	0	94	0
21	2	52	94	0
22	2	32	95	0
23	2	52	94	0
24	2	0	94	0
25	0	0	94	0
27	2	52	95	0
28	0	0	95	0
29	2	52	94	0
30	2	15	95	0
31	2	52	95	0
..
74	2	52	95	0
75	0	0	95	0
76	2	52	95	0
77	2	52	95	0
78	2	52	95	0
79	2	0	94	0
80	2	52	95	1
81	0	0	94	0
82	2	45	95	0
83	2	52	95	0
84	0	0	94	0
85	2	43	94	0
86	2	4	94	0
88	2	0	95	1
89	2	52	94	0
90	2	52	95	0
91	0	0	95	0
93	2	52	94	0
94	2	0	94	0
95	0	0	94	0
96	2	52	94	0
98	0	0	95	0
99	2	52	95	0
100	2	0	95	0
101	2	40	94	0
102	0	0	94	0
103	2	50	94	0
104	2	0	94	0
105	0	0	94	0
106	2	4	95	0

[100 rows x 30 columns]

```
In [453]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 145005 entries, 0 to 199520
```

Data columns (total 30 columns):

```
Age                145005 non-null int64
ClassOfWorker      145005 non-null object
IndustryCode       145005 non-null int64
OccupationCode     145005 non-null int64
Education          145005 non-null object
WagePerHour        145005 non-null int64
MaritalStatus      145005 non-null object
MajorIndustryCode  145005 non-null object
MajorOccupationCode 145005 non-null object
Race               145005 non-null object
HispanicOrigin     145005 non-null object
Sex                145005 non-null object
FullOrPartTime     145005 non-null object
CapitalGains       145005 non-null int64
CapitalLosses      145005 non-null int64
StockDividends     145005 non-null int64
TaxFilerStat       145005 non-null object
HouseholdFamilyStatus 145005 non-null object
HouseholdSummary   145005 non-null object
LiveInHouse1Y      145005 non-null object
NumPersonsWorkedEmployer 145005 non-null int64
CountryBirthFather  145005 non-null object
CountryBirthMother  145005 non-null object
CountryBirthSelf    145005 non-null object
Citizenship         145005 non-null object
OwnBusiness         145005 non-null int64
VeteranBenefits     145005 non-null int64
WeeksWorkedInY      145005 non-null int64
Year                145005 non-null int64
Income>50k          145005 non-null int32
dtypes: int32(1), int64(12), object(17)
memory usage: 33.7+ MB
```

In [454]: df.describe()

```
Out[454]:
```

	Age	IndustryCode	OccupationCode	WagePerHour	\
count	145005.000000	145005.000000	145005.000000	145005.000000	
mean	39.418468	20.221965	14.959374	73.836302	
std	19.327772	18.201672	14.910480	313.499755	
min	0.000000	0.000000	0.000000	0.000000	
25%	25.000000	0.000000	0.000000	0.000000	
50%	38.000000	24.000000	12.000000	0.000000	
75%	52.000000	37.000000	29.000000	0.000000	
max	90.000000	51.000000	46.000000	9999.000000	

	CapitalGains	CapitalLosses	StockDividends	NumPersonsWorkedEmployer	\
--	--------------	---------------	----------------	--------------------------	---

count	145005.000000	145005.000000	145005.000000	145005.000000
mean	557.842136	48.684659	252.061688	2.581587
std	5286.552736	308.852831	2260.621084	2.402695
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	2.000000
75%	0.000000	0.000000	0.000000	5.000000
max	99999.000000	4608.000000	99999.000000	6.000000

	OwnBusiness	VeteranBenefits	WeeksWorkedInY	Year \
count	145005.000000	145005.000000	145005.000000	145005.000000
mean	0.234282	1.826165	30.532313	94.488976
std	0.629074	0.551837	23.667317	0.499880
min	0.000000	0.000000	0.000000	94.000000
25%	0.000000	2.000000	0.000000	94.000000
50%	0.000000	2.000000	47.000000	94.000000
75%	0.000000	2.000000	52.000000	95.000000
max	2.000000	2.000000	52.000000	95.000000

	Income>50k
count	145005.000000
mean	0.080577
std	0.272185
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

Looks like there are some wrong values, such as WagePerHour (9999). I will look into those during EDA

4 Exploratory Data Analysis

```
In [455]: df['Income>50k'].value_counts()
```

```
Out[455]: 0    133321
          1     11684
          Name: Income>50k, dtype: int64
```

```
In [456]: df['Income>50k'].value_counts(normalize=True)
```


```
Out[456]: 0    0.919423
          1    0.080577
          Name: Income>50k, dtype: float64
```

Age

```
In [457]: df['Age'].describe()
```


```
Out [457]: count    145005.000000
           mean      39.418468
           std       19.327772
           min        0.000000
           25%       25.000000
           50%       38.000000
           75%       52.000000
           max       90.000000
           Name: Age, dtype: float64
```

```
In [458]: sns.distplot(df['Age'], bins=10)
           plt.grid()
```



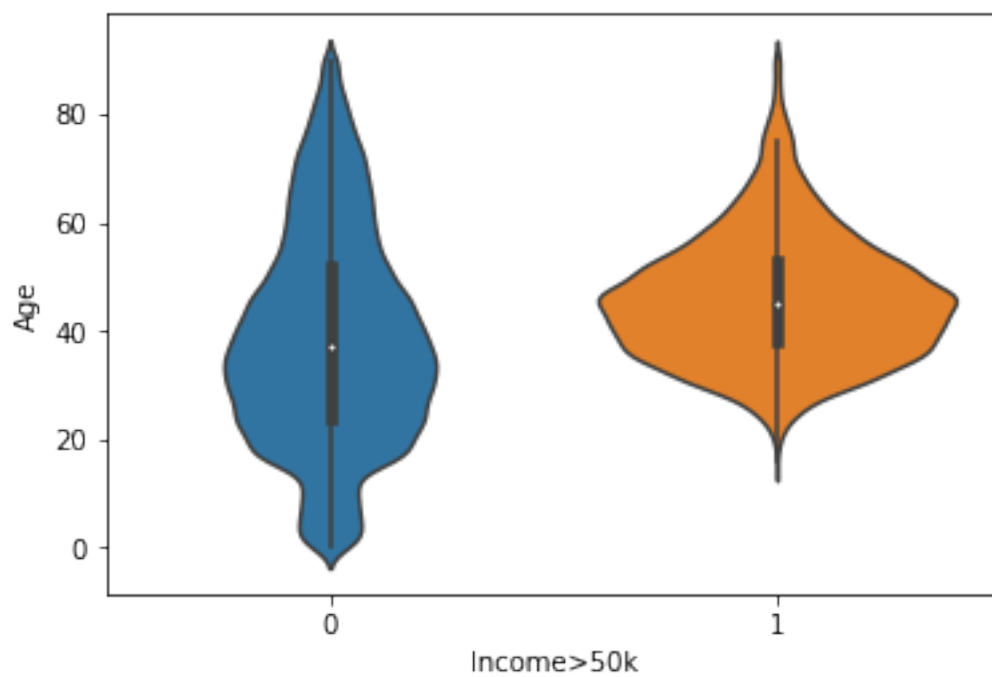
```
In [459]: sns.boxplot(x='Income>50k', y='Age', data=df)
```

```
Out [459]: <matplotlib.axes._subplots.AxesSubplot at 0x12cc1a08160>
```



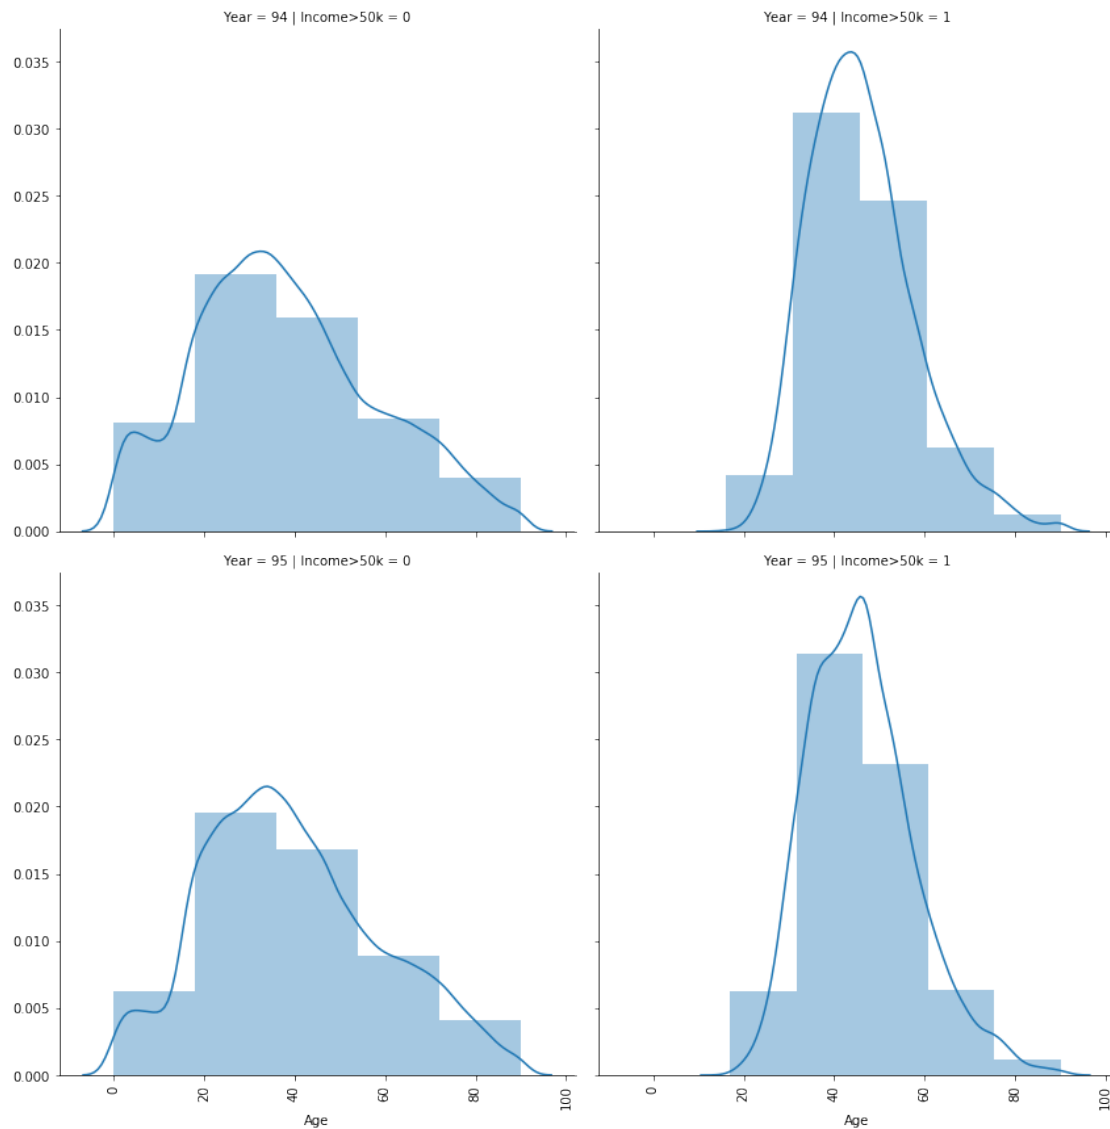
```
In [460]: sns.violinplot(x='Income>50k', y='Age', data=df)
```

```
Out[460]: <matplotlib.axes._subplots.AxesSubplot at 0x12da2f2cb00>
```



```
In [461]: plt.figure(figsize=(10,10))
g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.distplot, 'Age', bins=5)
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```

<Figure size 720x720 with 0 Axes>



ClassOfWorker



```
In [462]: df['ClassOfWorkers'].describe()
```

```
Out[462]: count      145005
unique          9
top      Private
freq      69208
Name: ClassOfWorkers, dtype: object
```

```
In [463]: df['ClassOfWorkers'].value_counts()
```

```
Out[463]: Private      69208
NA      49901
Self-employed-not incorporated      7935
Local government      7512
State government      4090
Self-employed-incorporated      3033
Federal government      2813
Never worked      357
Without pay      156
Name: ClassOfWorkers, dtype: int64
```

```
In [464]: plt.figure(figsize=(12,8))
sns.barplot(y='ClassOfWorkers', x='Income>50k', data=df, orient="h", hue='Year', dodge=True)
plt.grid(True)
```



```
In [465]: #Rename Never Worked and Without Pay to NA
          #Join State and Local Govt
```

```

df['ClassOfWorkers'] = np.where(df['ClassOfWorkers'] == 'Never worked', 'NA', df['Class
df['ClassOfWorkers'] = np.where(df['ClassOfWorkers'] == 'Without pay', 'NA', df['Class
df['ClassOfWorkers'] = np.where(df['ClassOfWorkers'] == 'Local government', 'Non Feder
df['ClassOfWorkers'] = np.where(df['ClassOfWorkers'] == 'State government', 'Non Feder

df_test['ClassOfWorkers'] = np.where(df_test['ClassOfWorkers'] == 'Never worked', 'NA'
df_test['ClassOfWorkers'] = np.where(df_test['ClassOfWorkers'] == 'Without pay', 'NA',
df_test['ClassOfWorkers'] = np.where(df_test['ClassOfWorkers'] == 'Local government',
df_test['ClassOfWorkers'] = np.where(df_test['ClassOfWorkers'] == 'State government',

df['ClassOfWorkers'].value_counts()

```

```

Out[465]: Private                69208
        NA                    50414
        Non Federal Government  11602
        Self-employed-not incorporated    7935
        Self-employed-incorporated    3033
        Federal government    2813
        Name: ClassOfWorkers, dtype: int64

```

```


In [466]: sns.boxplot(y='ClassOfWorkers', x='CapitalGains', data=df, orient='h')

```

```

Out[466]: <matplotlib.axes._subplots.AxesSubplot at 0x12d9b2f4160>

```



Industry Code - Major Industry Code

```

In [467]: df['IndustryCode'].unique()

```

```

Out[467]: array([ 0,  4, 40, 34, 43, 37, 24, 39, 12, 35, 45,  3, 19, 29, 32, 48, 33,
        23, 44, 36, 31, 30, 41,  5, 11,  9, 42,  6, 18, 50,  2,  1, 26, 47,
        16, 14, 22, 17,  7,  8, 25, 46, 27, 15, 13, 49, 38, 21, 28, 51, 20,
        10], dtype=int64)

```

```
In [468]: df['IndustryCode'].value_counts()
```

```
Out[468]: 0      50258
          33      16241
          43       7922
           4       5771
          42      4493
          45      4226
          29      4041
          37      3817
          41      3808
          32      3458
          35      3254
          39      2797
          34      2629
          44      2457
           2      2072
          11      1695
          50      1656
          47      1609
          40      1589
          38      1544
          24      1444
          19      1319
          12      1302
          31      1143
          30      1141
          25      1023
           9       967
          22       917
          36       892
          13       874
           1       808
          48       626
          27       609
          49       587
           3       557
           5       548
          21       540
           6       539
           8       535
          16       516
          23       516
          18       461
          15       437
           7       412
          14       289
          46       183
```

```

17      153
28      138
26      124
51       33
20       31
10        4
Name: IndustryCode, dtype: int64

```

```


In [469]: plt.figure(figsize=(20,20))
          sns.countplot(y='Income>50k', hue='IndustryCode', data=df[df['Income>50k']==1], palette=
          plt.grid()
          plt.legend(fontsize='large')

```


```

Out[469]: <matplotlib.legend.Legend at 0x12d9b24d630>

```



```
In [470]: plt.figure(figsize=(12,20))
          sns.barplot(y='IndustryCode', x='Income>50k', data=df, orient="h", dodge=True)
          plt.grid(True)
```




```
In [471]: df['MajorIndustryCode'].value_counts()
```


```
Out[471]: Not in universe or children      50258
          Retail trade                     16241
          Manufacturing-durable goods      8732
          Education                       7922
          Manufacturing-nondurable goods   6661
          Finance insurance and real estate 5883
          Construction                    5771
          Business and repair services     5361
          Medical except hospital          4493
          Public administration            4478
          Other professional services      4226
          Transportation                   4041
          Hospital services                3808
          Wholesale trade                  3458
          Agriculture                     2880
          Personal services except private HH 2797
          Social services                  2457
          Entertainment                   1589
          Utilities and sanitary services  1143
          Communications                   1141
          Private household services       892
          Mining                          557
          Forestry and fisheries           183
          Armed Forces                     33
          Name: MajorIndustryCode, dtype: int64
```

```
In [472]: plt.figure(figsize=(15,10))
          sns.countplot(y='Income>50k', hue='MajorIndustryCode', data=df[df['Income>50k']==1],
          plt.grid()
          plt.legend(fontsize='medium')
```

```
Out[472]: <matplotlib.legend.Legend at 0x12d9a1b8358>
```



```
In [473]: plt.figure(figsize=(12,30))
sns.barplot(y='MajorIndustryCode', x='Income>50k', data=df, orient="h", hue='Year')
plt.grid(True)
```

```
In [474]: df.groupby('MajorIndustryCode')['IndustryCode'].unique()
```

```
Out[474]: MajorIndustryCode
Agriculture
Armed Forces
Business and repair services
Communications
Construction
Education
Entertainment
Finance insurance and real estate
Forestry and fisheries
Hospital services
Manufacturing-durable goods
Manufacturing-nondurable goods
Medical except hospital
Mining
Not in universe or children
Other professional services
Personal services except private HH
Private household services
Public administration
Retail trade
Social services
Transportation
Utilities and sanitary services
Wholesale trade
Name: IndustryCode, dtype: object
```

```
In [475]: df.groupby(['IndustryCode', 'MajorIndustryCode'])['Income>50k'].sum().sort_values()
```

```
Out[475]: IndustryCode MajorIndustryCode
10 Manufacturing-durable goods 0
36 Private household services 4
28 Manufacturing-nondurable goods 7
51 Armed Forces 7
17 Manufacturing-durable goods 12
20 Manufacturing-nondurable goods 15
46 Forestry and fisheries 25
6 Manufacturing-durable goods 26
21 Manufacturing-nondurable goods 26
22 Manufacturing-nondurable goods 32
18 Manufacturing-durable goods 35
5 Manufacturing-durable goods 36
7 Manufacturing-durable goods 44
1 Agriculture 45
```

26	Manufacturing-nondurable goods	48
27	Manufacturing-nondurable goods	56
48	Public administration	63
14	Manufacturing-durable goods	69
44	Social services	79
8	Manufacturing-durable goods	86
19	Manufacturing-nondurable goods	90
38	Business and repair services	90
23	Manufacturing-nondurable goods	92
39	Personal services except private HH	98
40	Entertainment	103
15	Manufacturing-durable goods	109
16	Manufacturing-durable goods	115
9	Manufacturing-durable goods	121
2	Agriculture	127
3	Mining	143
49	Public administration	148
24	Manufacturing-nondurable goods	167
13	Manufacturing-durable goods	202
31	Utilities and sanitary services	245
30	Communications	262
12	Manufacturing-durable goods	267
25	Manufacturing-nondurable goods	276
50	Public administration	282
47	Public administration	308
11	Manufacturing-durable goods	310
29	Transportation	453
41	Hospital services	455
37	Business and repair services	482
34	Finance insurance and real estate	494
4	Construction	508
32	Wholesale trade	524
42	Medical except hospital	557
35	Finance insurance and real estate	568
33	Retail trade	750
0	Not in universe or children	820
43	Education	840
45	Other professional services	963

Name: Income>50k, dtype: int32

```
In [476]: #Industry Code is too scattered, and the values are too small for each Code. I will
df.drop('IndustryCode', axis=1, inplace=True)
df_test.drop('IndustryCode', axis=1, inplace=True)
```

OccupationCode and MajorOccupationCode

```
In [477]: df['OccupationCode'].describe()
```

```
Out[477]: count      145005.000000
mean              14.959374
```

```
std          14.910480
min           0.000000
25%           0.000000
50%          12.000000
75%          29.000000
max          46.000000
Name: OccupationCode, dtype: float64
```

```
In [478]: df['OccupationCode'].value_counts()
```


```
Out[478]: 0      50258
          2       8368
          26      7613
          19      5130
          29      4904
          36      4011
          34      3855
          10      3531
          23      3283
          16      3272
          33      3200
          12      3171
          35      3064
           3      3044
          38      2894
          31      2601
          32      2295
          37      2169
           8      2063
          42      1873
          30      1811
          24      1784
          17      1700
          28      1608
          44      1553
          41      1525
          43      1273
           4      1269
          13      1230
          18      1035
          39      1001
          14       909
           5       806
          15       777
          25       745
          27       733
           9       681
           7       663
```

```
40      604
11      591
1       519
21      518
6       404
22      397
45      169
20       68
46       33
Name: OccupationCode, dtype: int64
```

```
In [479]: df['MajorOccupationCode'].unique()
```

```
Out[479]: array(['NA', 'Precision production craft & repair',
                'Professional specialty', 'Executive admin and managerial',
                'Handlers equip cleaners etc', 'Adm support including clerical',
                'Machine operators assmblrs & inspctrs', 'Other service', 'Sales',
                'Private household services', 'Technicians and related support',
                'Transportation and material moving',
                'Farming forestry and fishing', 'Protective services',
                'Armed Forces'], dtype=object)
```


```
In [480]: plt.figure(figsize=(12,20))
          sns.barplot(y='OccupationCode', x='Income>50k', data=df, orient="h", hue='Year', dodge=True)
          plt.grid(True)
```



```
In [481]: df.groupby('MajorOccupationCode')['OccupationCode'].unique()
```


```
Out[481]: MajorOccupationCode
Adm support including clerical      [26, 22, 25, 23, 21, 24]
Armed Forces                        [46]
Executive admin and managerial      [3, 2, 1]
Farming forestry and fishing        [44, 43, 45]
Handlers equip cleaners etc         [40, 41, 42]
Machine operators assmblrs & inspctrs [37, 36]
NA                                  [0]
Other service                       [31, 29, 32, 30]
Precision production craft & repair  [34, 35, 33]
Private household services          [27]
Professional specialty              [10, 12, 8, 9, 11, 7, 4, 6, 5]
Protective services                 [28]
Sales                              [19, 16, 18, 17, 20]
Technicians and related support      [13, 14, 15]
Transportation and material moving   [39, 38]
Name: OccupationCode, dtype: object
```

```
In [482]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'OccupationCode', order=df['OccupationCode'].unique())
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=60) # set new labels
```




```
In [483]: plt.figure(figsize=(15,15))
          sns.countplot(y='Income>50k', hue='OccupationCode', data=df[df['Income>50k']==1], palette='magma',
          plt.grid()
          plt.legend(fontsize='medium')
```

```
Out[483]: <matplotlib.legend.Legend at 0x12d99f609e8>
```





```
In [484]: plt.figure(figsize=(12,20))
          sns.barplot(y='MajorOccupationCode', x='Income>50k', data=df, orient="h", hue='Year')
          plt.grid(True)
```



```
In [485]: plt.figure(figsize=(15,10))
sns.countplot(y='Income>50k', hue='MajorOccupationCode', data=df[df['Income>50k']==1])
plt.grid()
plt.legend(fontsize='medium')
```

```
Out[485]: <matplotlib.legend.Legend at 0x12c80222ef0>
```



```
In [486]: #Occupation Code is too scattered, and the values are too small for each Code. I will
df.drop('OccupationCode', axis=1, inplace=True)
df_test.drop('OccupationCode', axis=1, inplace=True)
```





Education

```
In [487]: df['Education'].value_counts()
```

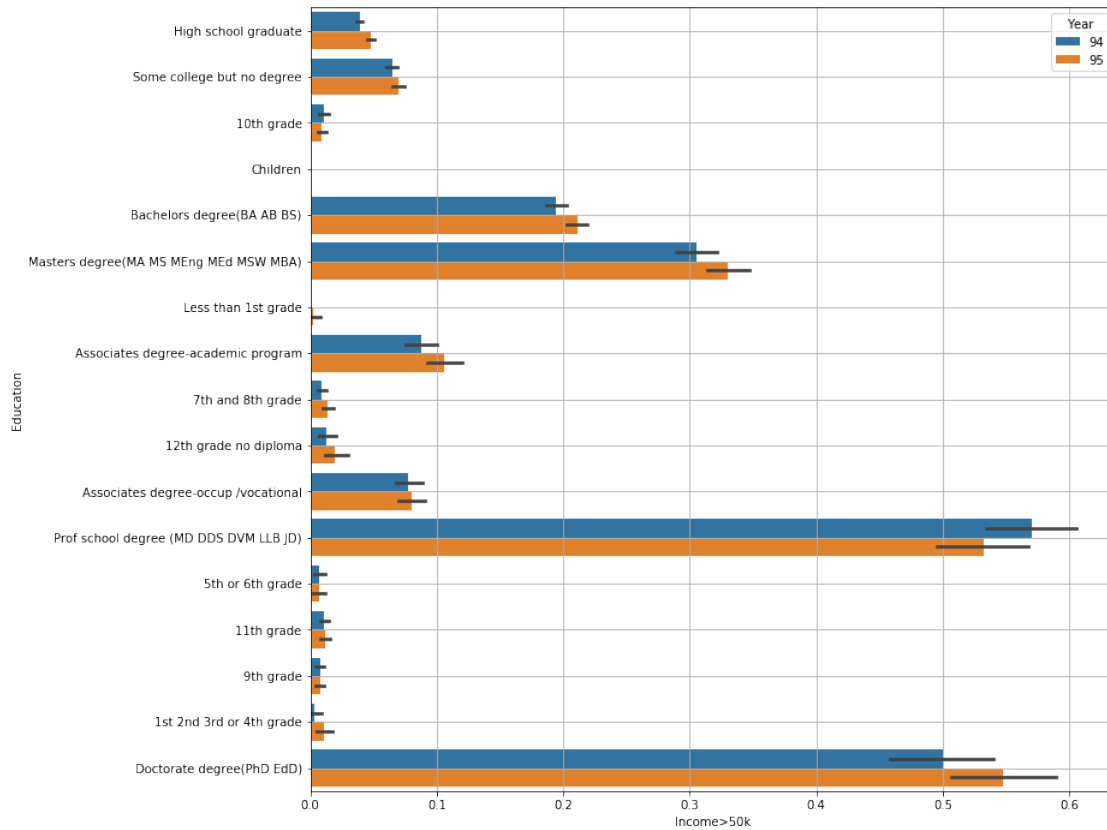
```
Out[487]: High school graduate          41733
Some college but no degree          25146
Bachelors degree(BA AB BS)         18278
Children                            11679
10th grade                          6199
Masters degree(MA MS MEng MEd MSW MBA) 6059
11th grade                          6003
7th and 8th grade                   5852
Associates degree-occup /vocational 5034
9th grade                          4787
Associates degree-academic program  4125
5th or 6th grade                    2999
12th grade no diploma               1949
1st 2nd 3rd or 4th grade            1683
Prof school degree (MD DDS DVM LLB JD) 1598
Doctorate degree(PhD EdD)           1120
```

```
Less than 1st grade                                761
Name: Education, dtype: int64
```

```
In [488]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'Education', order=df['Education'].unique() )
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```




```
In [489]: plt.figure(figsize=(12,12))
sns.barplot(y='Education', x='Income>50k', data=df, orient="h", hue='Year', dodge=True)
plt.grid(True)
```



```
In [490]: plt.figure(figsize=(10,10))
sns.countplot(y='Income>50k', hue='Education', data=df[df['Income>50k']==1],palette=
plt.grid()
plt.legend(fontsize='medium')
```

```
Out[490]: <matplotlib.legend.Legend at 0x12cd3999588>
```



```
In [491]: df['Education'].unique()
```

```
Out[491]: array(['High school graduate', 'Some college but no degree', '10th grade',
                  'Children', 'Bachelors degree(BA AB BS)',
                  'Masters degree(MA MS MEng MEd MSW MBA)', 'Less than 1st grade',
                  'Associates degree-academic program', '7th and 8th grade',
                  '12th grade no diploma', 'Associates degree-occup /vocational',
                  'Prof school degree (MD DDS DVM LLB JD)', '5th or 6th grade',
                  '11th grade', '9th grade', '1st 2nd 3rd or 4th grade',
                  'Doctorate degree(PhD EdD)'], dtype=object)
```

```
In [492]: #I will group all Education levels other than 'High school graduate', 'Some college but no degree',
           #Bachelors degree(BA AB BS)' and 'Masters degree(MA MS MEng MEd MSW MBA)' as 'Other'
           #I will also rename 'Some college but no degree' as 'High school graduate'.
```

```

df['Education'] = df['Education'].apply(lambda x: 'Other' if x not in ['High school',
    'Bachelors degree(BA AB BS)', 'Masters degree(MA MS MEng MEd MSW MBA)',
    'Prof school degree (MD DDS DVM LLB JD)', 'Doctorate degree(PhD EdD)'] else x)

df['Education'] = np.where(df['Education'] == 'Some college but no degree', 'High school', df['Education'])

df_test['Education'] = df_test['Education'].apply(lambda x: 'Other' if x not in ['High school',
    'Bachelors degree(BA AB BS)', 'Masters degree(MA MS MEng MEd MSW MBA)',
    'Prof school degree (MD DDS DVM LLB JD)', 'Doctorate degree(PhD EdD)'] else x)

df_test['Education'] = np.where(df_test['Education'] == 'Some college but no degree', 'High school', df_test['Education'])

```

WagePerHour

```
In [493]: df.columns
```

```

Out[493]: Index(['Age', 'ClassOfWorker', 'Education', 'WagePerHour', 'MaritalStatus',
    'MajorIndustryCode', 'MajorOccupationCode', 'Race', 'HispanicOrigin',
    'Sex', 'FullOrPartTime', 'CapitalGains', 'CapitalLosses',
    'StockDividends', 'TaxFilerStat', 'HouseholdFamilyStatus',
    'HouseholdSummary', 'LiveInHouse1Y', 'NumPersonsWorkedEmployer',
    'CountryBirthFather', 'CountryBirthMother', 'CountryBirthSelf',
    'Citizenship', 'OwnBusiness', 'VeteranBenefits', 'WeeksWorkedInY',
    'Year', 'Income>50k'],
    dtype='object')


```

```

In [494]: plt.figure(figsize=(10,10))
    sns.boxplot(x='Income>50k', y='WagePerHour', data=df)

```

```
Out[494]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4eb24d68>
```

```
In [495]: df.groupby('WagePerHour')['Age'].count()
```

```
Out[495]: WagePerHour
0         134017
20           1
70           1
75           2
100          11
110           1
125           1
135           1
143           1
150           6
170           1
173           1
```

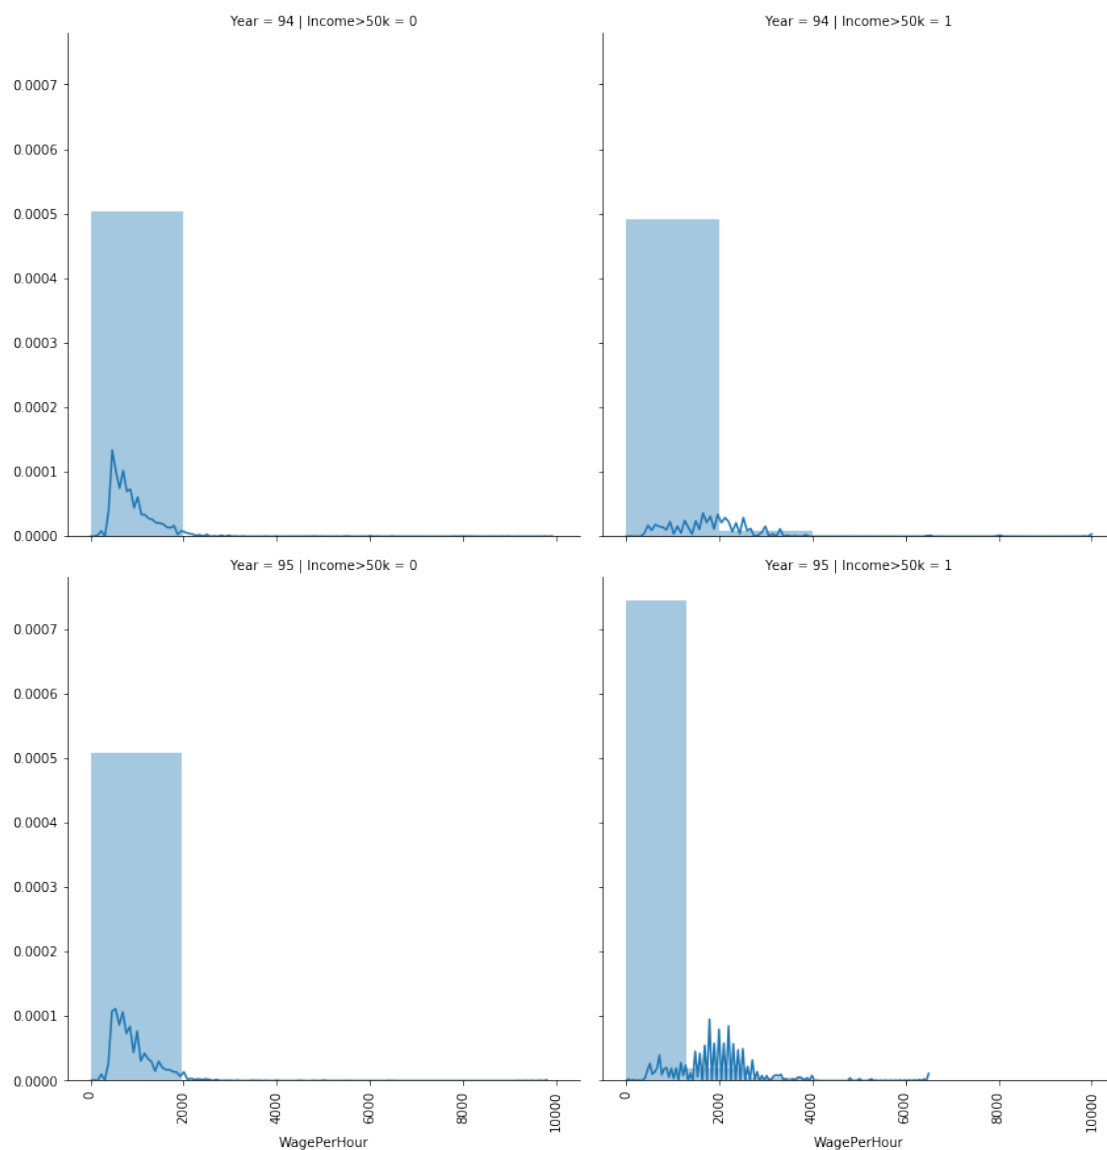
190	1
200	27
205	1
210	7
212	2
213	20
215	2
220	2
225	4
230	1
232	1
233	1
234	1
235	2
245	1
250	11
252	1
255	3
...	
4300	1
4500	4
4800	1
4807	5
4900	1
5000	5
5200	1
5250	1
5500	4
5525	1
6000	2
6009	1
6200	1
6410	1
6500	6
6600	1
6800	1
7000	1
7400	1
7700	1
7800	1
8000	4
8300	1
8600	1
8800	1
9000	1
9400	1
9800	2
9916	1

```
9999          1
Name: Age, Length: 1227, dtype: int64
```

```
In [496]: plt.figure(figsize=(10,10))
```


```
g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.distplot, 'WagePerHour', bins=5)
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```

<Figure size 720x720 with 0 Axes>




```
In [497]: sns.distplot(df['WagePerHour'], bins=4)
```

```
Out[497]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4be8aeb8>
```



```
In [498]: sns.countplot(x='Income>50k', data=df[df['WagePerHour']>2000])
```

```
Out[498]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4be50eb8>
```



5 It looks like the WagePerHouse has some wrong inputs. I might delete the rows, but a big oercentage of them are positive


for Income>50K. I will cap the Wages at \$2000 per hour.

In [499]: *#I will cap the WagePerHour to \$2,000. Higher numbers do not seem to be correct.*

```
df['WagePerHour'] = np.where(df['WagePerHour'] > 2000, 2000, df['WagePerHour'])
df_test['WagePerHour'] = np.where(df_test['WagePerHour'] > 2000, 2000, df_test['WagePerHour'])
```

In [500]: `sns.countplot(x='Income>50k', data=df[df['WagePerHour']== 2000])`

Out[500]: `<matplotlib.axes._subplots.AxesSubplot at 0x12d4be1cf28>`



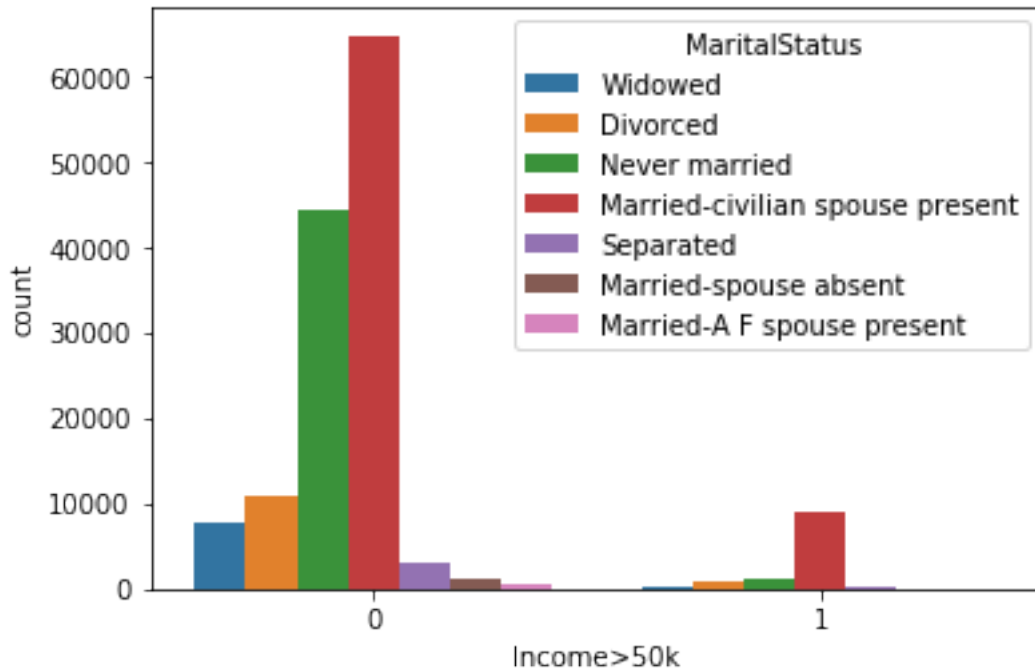
Marital Status

```
In [501]: df['MaritalStatus'].value_counts()
```

```
Out[501]: Married-civilian spouse present    73979  
Never married                             45594  
Divorced                                  12046  
Widowed                                   8032  
Separated                                 3325  
Married-spouse absent                     1396  
Married-A F spouse present                 633  
Name: MaritalStatus, dtype: int64
```


```
In [502]: sns.countplot(x='Income>50k', hue='MaritalStatus', data=df)
```

```
Out[502]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4bdfb0f0>
```




```
In [503]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'MaritalStatus', order=df['MaritalStatus'].unique())
g.set_xticklabels(rotation=30)

for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```




```
In [504]: plt.figure(figsize=(10,10))
          sns.countplot(x='Income>50k', hue='MaritalStatus', data=df[df['Income>50k']==1],pale
          plt.grid()
          plt.legend(fontsize='medium')
```

Out[504]: <matplotlib.legend.Legend at 0x12d4eb55320>



```
In [505]: plt.figure(figsize=(12,12))  
sns.barplot(y='MaritalStatus', x='Income>50k', data=df, orient="h", hue='Year', dodge=True)  
plt.grid(True)
```



```
In [506]: df['MaritalStatus'].unique()
```

```
Out[506]: array(['Widowed', 'Divorced', 'Never married',
                  'Married-civilian spouse present', 'Separated',
                  'Married-spouse absent', 'Married-A F spouse present'], dtype=object)
```

```
In [507]: #I will join some values under two: Married and Divorced
```

```
df['MaritalStatus'] = df['MaritalStatus'].apply(lambda x: 'Divorced' if x in ['Divorced', 'Separated']
                                                else 'Married' if x in ['Married-civilian spouse present', 'Married-spouse absent']
                                                else x)
```

```
df_test['MaritalStatus'] = df_test['MaritalStatus'].apply(lambda x: 'Divorced' if x in ['Divorced', 'Separated']
                                                           else 'Married' if x in ['Married-civilian spouse present', 'Married-spouse absent']
                                                           else x)
```

```
In [508]: df['MaritalStatus'].value_counts()
```

```
Out[508]: Married      76008
Never married    45594
Divorced        15371
```

```
Widowed          8032
Name: MaritalStatus, dtype: int64
```

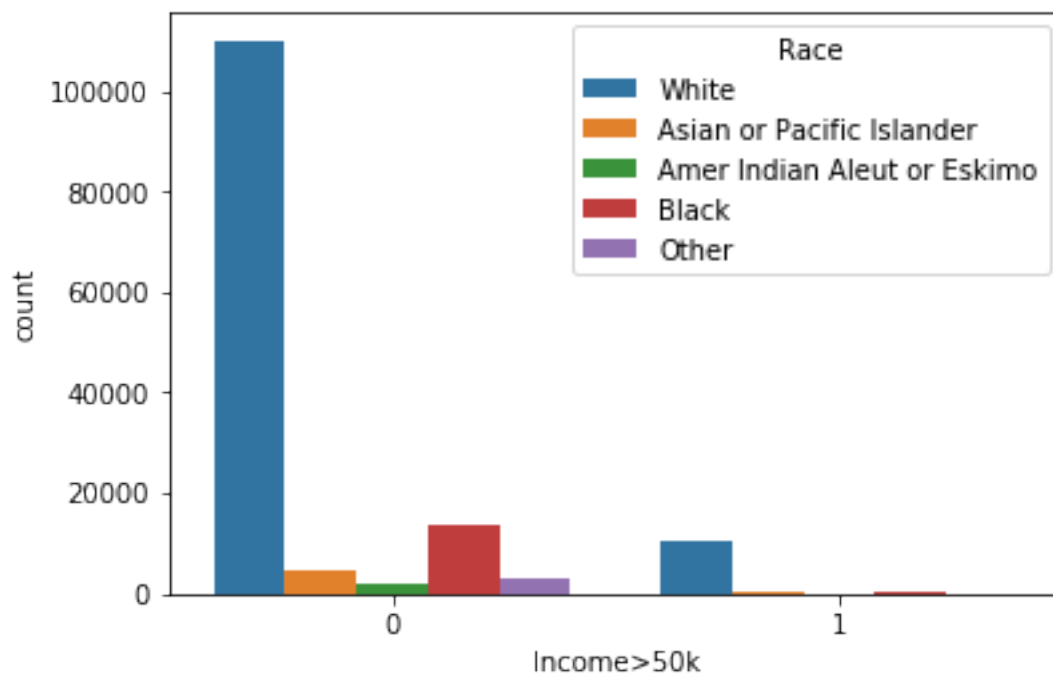
Race

```
In [509]: df['Race'].value_counts()
```

```
Out[509]: White          120679
Black          14424
Asian or Pacific Islander    4896
Other           3154
Amer Indian Aleut or Eskimo   1852
Name: Race, dtype: int64
```


```
In [510]: sns.countplot(x='Income>50k', hue='Race', data=df)
```

```
Out[510]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4ebc3c50>
```




```
In [511]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'Race', order=df['Race'].unique())
g.set_xticklabels(rotation=30)

for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```



```
In [512]: plt.figure(figsize=(12,12))
sns.barplot(y='Race', x='Income>50k', data=df, orient="h", hue='Year', dodge=True)
plt.grid(True)
```



```
In [513]: #I will add 'Amer Indian Aleut or Eskimo' to 'Other'
df['Race'] = df['Race'].apply(lambda x: 'Other' if x == 'Amer Indian Aleut or Eskimo')
df_test['Race'] = df_test['Race'].apply(lambda x: 'Other' if x == 'Amer Indian Aleut
```


HispanicOrigin

```
In [514]: df['HispanicOrigin'].value_counts(normalize=True)
```

```
Out[514]: All other                0.842978
Mexican (Mexicano)                0.043364
Mexican-American                  0.041185
Central or South American         0.021323
Puerto Rican                    0.020068
Other Spanish                     0.014606
Cuban                            0.006862
NA                                0.005593
Chicano                           0.002034
Do not know                       0.001986
Name: HispanicOrigin, dtype: float64
```

```
In [515]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'HispanicOrigin', order=df['HispanicOrigin'].unique())
g.set_xticklabels(rotation=30)
```

```
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```



```
In [516]: df['HispanicOrigin'].replace('All other', 'NA', inplace=True)
```


```

df_test['HispanicOrigin'].replace('All other', 'NA', inplace=True)


In [517]: g = sns.FacetGrid(data=df[df['HispanicOrigin'] != 'NA'], col='Income>50k', row='Year'
g.map(sns.countplot, 'HispanicOrigin', order=df['HispanicOrigin'].unique())
g.set_xticklabels(rotation=30)

for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels

```



```
In [518]: plt.figure(figsize=(12,12))
sns.barplot(y='HispanicOrigin', x='Income>50k', data=df[df['HispanicOrigin'] != 'NA'])
plt.grid(True)
```




```
In [519]: #I will drop this column. I will only be relying on Race
df.drop('HispanicOrigin', axis=1, inplace=True)
df_test.drop('HispanicOrigin', axis=1, inplace=True)
```


Sex

```
In [520]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'Sex', order=df['Sex'].unique())

for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```

```
In [521]: plt.figure(figsize=(12,12))
          sns.barplot(x='Sex', y='Income>50k', data=df, dodge=True)
          plt.grid(True)
```



```
In [522]: df.rename(columns={'Sex':'Male'}, inplace=True)
df['Male'] = np.where(df['Male']=='Male',1,0)
df_test.rename(columns={'Sex':'Male'}, inplace=True)
df_test['Male'] = np.where(df_test['Male']=='Male',1,0)
```





FullOrPartTime

```
In [523]: df['FullOrPartTime'].value_counts()
```


```
Out[523]: Children or Armed Forces      78981
Full-time schedules                    38956
Not in labor force                     19201
PT for non-econ reasons usually FT    3210
Unemployed full-time                  2200
PT for econ reasons usually PT        1166
```

```
Unemployed part- time          779
PT for econ reasons usually FT  512
Name: FullOrPartTime, dtype: int64
```

```
In [524]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'FullOrPartTime', order=df['FullOrPartTime'].unique() )
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```




```
In [525]: plt.figure(figsize=(12,12))
sns.barplot(x='FullOrPartTime', y='Income>50k', data=df, hue='Year', dodge=True)
plt.xticks(rotation=45)
plt.grid(True)
```



Looks like the data for this column is missing for 1994. It doesn't make sense that all those who answered the survey in 1994 are of the same type


```
In [526]: sns.countplot(x='Income>50k', hue='FullOrPartTime', data=df[df['Year']==94])
```

```
Out[526]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4fb44ac8>
```



```
In [527]: plt.figure(figsize=(10,10))
          sns.countplot(x='Income>50k', hue='FullOrPartTime', data=df[df['Income>50k']==1], palette='magma',
          plt.grid()
          plt.legend(fontsize='medium')
```

```
Out[527]: <matplotlib.legend.Legend at 0x12d4fb46da0>
```



Even though this columns seems to be correlated with Income levels, but I will have to drop it as I cannot impute the values of 1994

```
In [528]: df.drop('FullOrPartTime', axis=1, inplace=True)
          df_test.drop('FullOrPartTime', axis=1, inplace=True)
```

CapitalGains & CapitalLosses

```
In [529]: df['CapitalGains'].describe()
```

```
Out[529]: count    145005.000000
          mean       557.842136
          std       5286.552736
          min         0.000000
          25%         0.000000
          50%         0.000000
```

```

75%          0.000000
max          99999.000000
Name: CapitalGains, dtype: float64

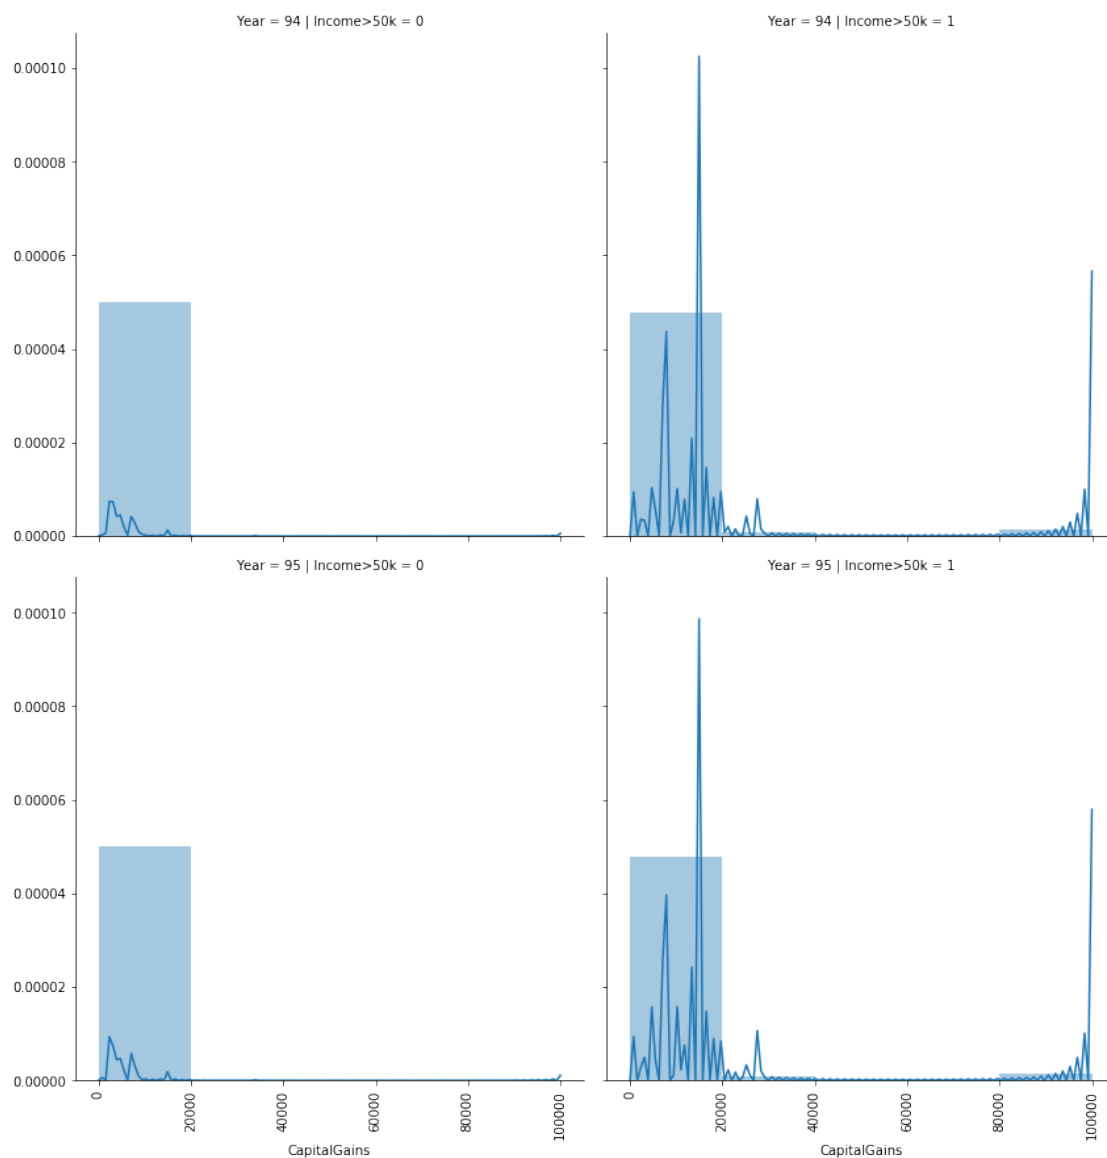
```

```

In [530]: plt.figure(figsize=(10,10))
g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.distplot, 'CapitalGains', bins=5)
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels


```

<Figure size 720x720 with 0 Axes>





```
In [531]: sns.distplot(df['CapitalGains'], bins=4)
```

```
Out[531]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4f0c29b0>
```



```
In [532]: plt.figure(figsize=(8,8))
          sns.boxplot(y='CapitalGains', x='Income>50k', data=df[df['CapitalGains']>0])
```

```
Out[532]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4f0af390>
```



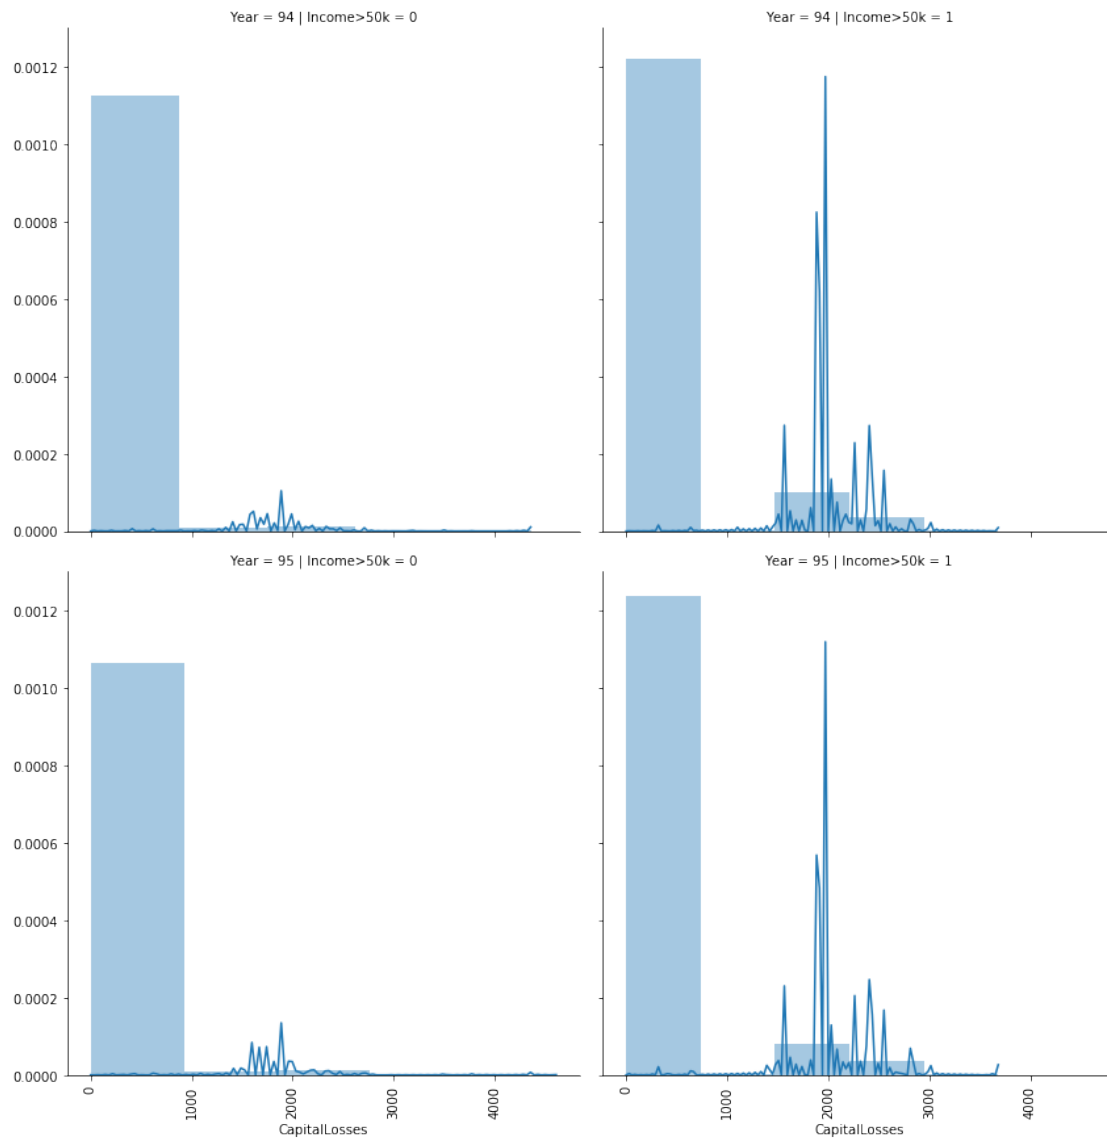
```
In [533]: df['CapitalLosses'].describe()
```

```
Out[533]: count      145005.000000
          mean         48.684659
          std         308.852831
          min           0.000000
          25%           0.000000
          50%           0.000000
          75%           0.000000
          max         4608.000000
          Name: CapitalLosses, dtype: float64
```

```
In [534]: plt.figure(figsize=(10,10))
          g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
          g.map(sns.distplot, 'CapitalLosses', bins=5)
          for ax in g.axes.flat:
```


```
labels = ax.get_xticklabels() # get x labels
ax.set_xticklabels(labels, rotation=90) # set new labels
```

<Figure size 720x720 with 0 Axes>




```
In [535]: sns.distplot(df['CapitalLosses'], bins=4)
```

```
Out[535]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4ebc07f0>
```




```
In [536]: plt.figure(figsize=(8,8))  
          sns.boxplot(y='CapitalLosses', x='Income>50k', data=df[df['CapitalLosses']>0])  
  
Out[536]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4ebcc860>
```




```
In [537]: plt.figure(figsize=(8,8))
          sns.violinplot(y='CapitalLosses', x='Income>50k', data=df[df['CapitalLosses']>0])

Out[537]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4ebc42e8>
```



```
In [538]: plt.figure(figsize=(8,8))
          sns.swarmplot(y='CapitalLosses', x='Income>50k', data=df[df['CapitalLosses']>0])

Out[538]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4eb964a8>
```



This is unexpected, since higher income should probably mean lower Capital Losses.
StockDividends

```
In [539]: df['StockDividends'].describe()
```

```
Out[539]: count      145005.000000
          mean         252.061688
          std        2260.621084
          min           0.000000
          25%           0.000000
          50%           0.000000
          75%           0.000000
          max        99999.000000
          Name: StockDividends, dtype: float64
```

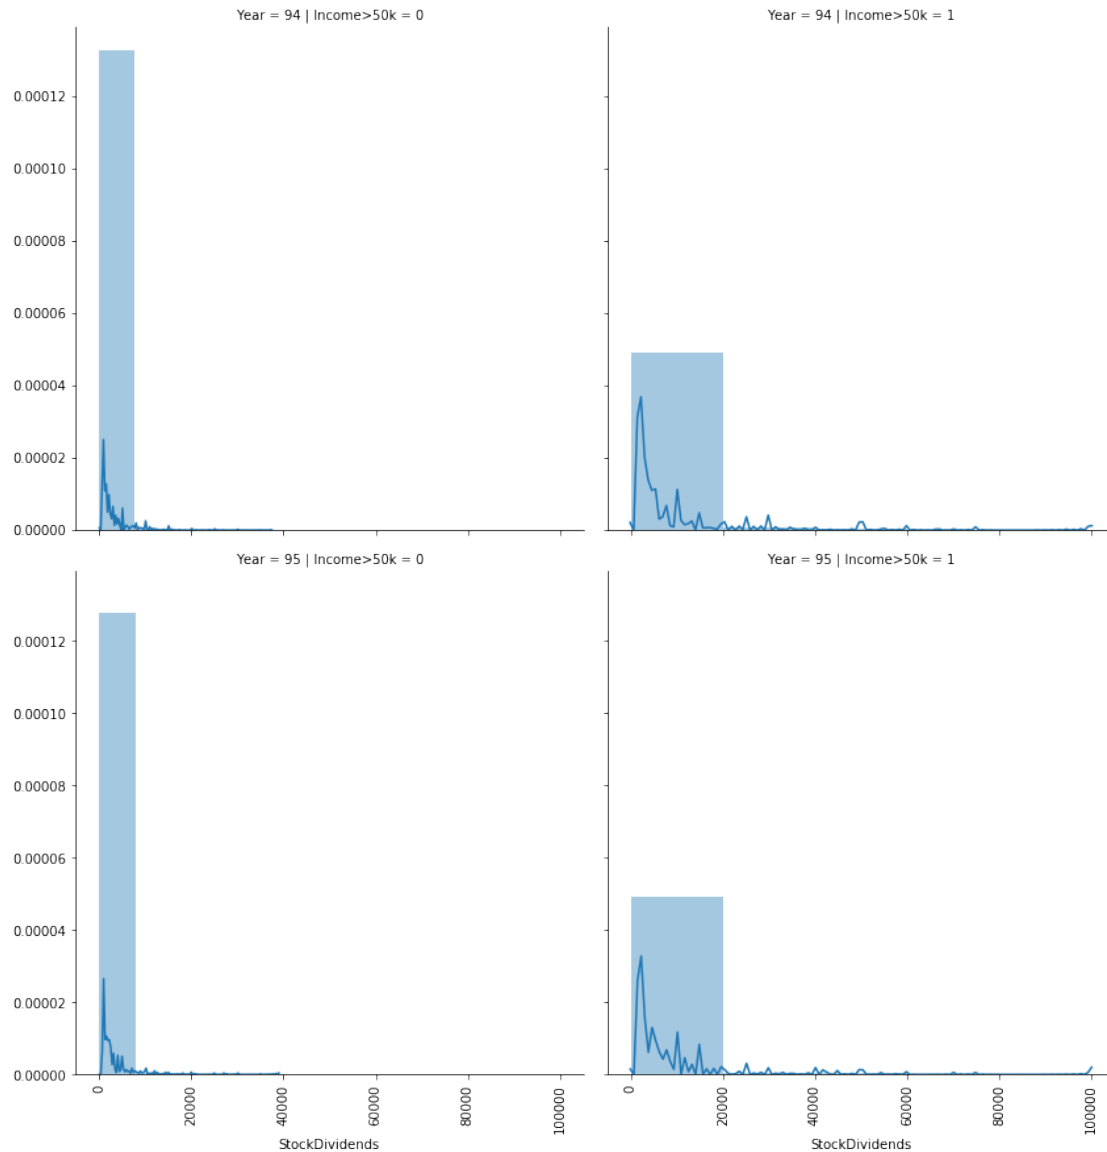
```
In [540]: plt.figure(figsize=(10,10))
          g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
```

```

g.map(sns.distplot, 'StockDividends', bins=5)
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels

```

<Figure size 720x720 with 0 Axes>




```

In [541]: sns.distplot(df['StockDividends'], bins=4)
plt.xlim(0,30000)


```

Out [541]: (0, 30000)



```
In [542]: plt.figure(figsize=(8,8))
          sns.boxplot(y='StockDividends', x='Income>50k', data=df[df['StockDividends']>0])

Out[542]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4fb36cf8>
```



TaxFilerStat

```
In [543]: df['TaxFilerStat'].describe()
```

```
Out[543]: count          145005
          unique           6
          top      Joint both under 65
          freq          61723
          Name: TaxFilerStat, dtype: object
```

```
In [544]: df['TaxFilerStat'].value_counts()
```

```
Out[544]: Joint both under 65      61723
          Single                  34850
          Nonfiler                 32260
          Head of household         7195
          Joint both 65+           5767
```

```


Joint one under 65 & one 65+      3210
Name: TaxFilerStat, dtype: int64

```


```

In [545]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'TaxFilerStat', order=df['TaxFilerStat'].unique() )
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels

```




```
In [546]: plt.figure(figsize=(12,12))
sns.barplot(x='TaxFilerStat', y='Income>50k', data=df, hue='Year', dodge=True)
plt.xticks(rotation=45)
plt.grid(True)
```



```
In [547]: plt.figure(figsize=(10,10))
sns.countplot(x='Income>50k', hue='TaxFilerStat', data=df[df['Income>50k']==1], palette='magma')
plt.grid()
plt.legend(fontsize='medium')
```

Out [547]: <matplotlib.legend.Legend at 0x12d4fb514e0>



```
In [548]: #Group the values: ['Head of household', 'Joint both 65+', 'Joint one under 65 & one 65+']
df['TaxFilerStat'] = df['TaxFilerStat'].apply(lambda x: 'Other' if x in ['Head of household', 'Joint both 65+', 'Joint one under 65 & one 65+'] else x)
df_test['TaxFilerStat'] = df_test['TaxFilerStat'].apply(lambda x: 'Other' if x in ['Head of household', 'Joint both 65+', 'Joint one under 65 & one 65+'] else x)
```

HouseholdFamilyStatus


```
In [549]: df['HouseholdFamilyStatus'].describe()
```

```
Out [549]: count      145005
unique         38
top      Householder
freq         48707
Name: HouseholdFamilyStatus, dtype: object
```

```
In [550]: df['HouseholdFamilyStatus'].value_counts()
```


```
Out[550]: Householder                48707
          Spouse of householder       35347
          Nonfamily householder       19255
          Child <18 never marr not in subfamily 14015
          Child 18+ never marr Not in a subfamily 10969
          Secondary individual         5711
          Other Rel 18+ ever marr not in subfamily 1709
          Other Rel 18+ never marr not in subfamily 1611
          Child 18+ ever marr Not in a subfamily 991
          Grandchild <18 never marr child of subfamily RP 799
          RP of unrelated subfamily     663
          Child 18+ ever marr RP of subfamily 652
          Other Rel 18+ ever marr RP of subfamily 609
          Other Rel 18+ spouse of subfamily RP 591
          Child 18+ never marr RP of subfamily 569
          Other Rel <18 never marr child of subfamily RP 515
          Grandchild <18 never marr not in subfamily 485
          Child under 18 of RP of unrel subfamily 412
          Other Rel <18 never marr not in subfamily 391
          Grandchild 18+ never marr not in subfamily 364
          In group quarters             179
          Child 18+ spouse of subfamily RP 121
          Other Rel 18+ never marr RP of subfamily 92
          Child <18 never marr RP of subfamily 76
          Spouse of RP of unrelated subfamily 49
          Child <18 ever marr not in subfamily 35
          Grandchild 18+ ever marr not in subfamily 34
          Grandchild 18+ spouse of subfamily RP 10
          Grandchild 18+ ever marr RP of subfamily 9
          Child <18 ever marr RP of subfamily 9
          Grandchild 18+ never marr RP of subfamily 6
          Other Rel <18 ever marr RP of subfamily 6
          Other Rel <18 never married RP of subfamily 4
          Other Rel <18 spouse of subfamily RP 3
          Child <18 spouse of subfamily RP 2
          Grandchild <18 ever marr not in subfamily 2
          Grandchild <18 never marr RP of subfamily 2
          Other Rel <18 ever marr not in subfamily 1
          Name: HouseholdFamilyStatus, dtype: int64
```

```
In [551]: plt.figure(figsize=(12,12))
          sns.barplot(x='HouseholdFamilyStatus', y='Income>50k', data=df, dodge=True)
          plt.xticks(rotation=90)
          plt.grid(True)
```







```
In [552]: plt.figure(figsize=(10,10))
sns.countplot(x='Income>50k', hue='HouseholdFamilyStatus', data=df[df['Income>50k']!=0])
plt.grid()
plt.legend(fontsize='medium')
```

Out [552]: <matplotlib.legend.Legend at 0x12d4fb71dd8>



```
In [553]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'HouseholdFamilyStatus', order=df['HouseholdFamilyStatus'].unique())
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```

```
In [554]: #To simplify the model, I will replace all small values with Other
df['HouseholdFamilyStatus'] = df['HouseholdFamilyStatus'].apply(lambda x: 'Other' if
df_test['HouseholdFamilyStatus'] = df_test['HouseholdFamilyStatus'].apply(lambda x:
```

```
In [555]: df['HouseholdFamilyStatus'].value_counts()
```

```
Out[555]: Householder          48707
Other          41696
Spouse of householder    35347
Nonfamily householder    19255
Name: HouseholdFamilyStatus, dtype: int64
```

LiveInHouse1Y

```
In [556]: df['LiveInHouse1Y'].describe()
```

```
Out[556]: count          145005
unique              3
top    Not in universe under 1 year old
freq          71217
Name: LiveInHouse1Y, dtype: object
```

```
In [557]: df['LiveInHouse1Y'].value_counts()
```

```
Out[557]: Not in universe under 1 year old    71217
Yes          58972
No          14816
Name: LiveInHouse1Y, dtype: int64
```


```
In [558]: #Too much missing info, I will drop this column
df.drop('LiveInHouse1Y', axis=1, inplace=True)
df_test.drop('LiveInHouse1Y', axis=1, inplace=True)
```

NumPersonsWorkedEmployer


```
In [559]: df['NumPersonsWorkedEmployer'].describe()
```

```
Out[559]: count    145005.000000
mean           2.581587
std            2.402695
min            0.000000
25%            0.000000
50%            2.000000
75%            5.000000
max            6.000000
Name: NumPersonsWorkedEmployer, dtype: float64
```

```
In [560]: g = sns.FacetGrid(data=df, col='Income>50k', row='Year', height=6)
g.map(sns.countplot, 'NumPersonsWorkedEmployer', order=df['NumPersonsWorkedEmployer'].
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```




```
In [561]: plt.figure(figsize=(12,12))
sns.barplot(x='NumPersonsWorkedEmployer', y='Income>50k', data=df, dodge=True)
plt.xticks(rotation=90)
plt.grid(True)
```




```
In [562]: plt.figure(figsize=(10,10))
sns.countplot(x='Income>50k', hue='NumPersonsWorkedEmployer', data=df[df['Income>50k']
plt.grid()
plt.legend(fontsize='medium')
```

Out [562]: <matplotlib.legend.Legend at 0x12d9b33d358>



```
In [563]: sns.boxplot(x='Income>50k', y="NumPersonsWorkedEmployer", data=df)
```

```
Out[563]: <matplotlib.axes._subplots.AxesSubplot at 0x12d9b2e1438>
```




CountryBirthFather

```
In [564]: df['CountryBirthFather'].describe()
```

```
Out [564]: count          145005
           unique           42
           top      United-States
           freq          114596
           Name: CountryBirthFather, dtype: object
```


```
In [565]: plt.figure(figsize=(10,10))
           sns.countplot(y='Income>50k', hue='CountryBirthFather', data=df, orient='H',palette=
           plt.grid()
           plt.legend(fontsize='x-small')
```

```
Out [565]: <matplotlib.legend.Legend at 0x12d9b2a8ef0>
```




```
In [566]: plt.figure(figsize=(10,10))
          sns.countplot(y='Income>50k', hue='CountryBirthFather', data=df[df['Income>50k']==1])
          plt.grid()
          plt.legend(fontsize='x-small')
```


Out[566]: <matplotlib.legend.Legend at 0x12d9b36ce10>




```
In [567]: plt.figure(figsize=(12,12))
sns.barplot(x='CountryBirthFather', y='Income>50k', data=df, dodge=True)
plt.xticks(rotation=90)
plt.grid(True)
```

```
In [568]: plt.figure(figsize=(12,12))
sns.barplot(x='CountryBirthMother', y='Income>50k', data=df, dodge=True)
plt.xticks(rotation=90)
plt.grid(True)
```



```
In [569]: plt.figure(figsize=(12,12))
sns.barplot(x='CountryBirthSelf', y='Income>50k', data=df, dodge=True)
plt.xticks(rotation=90)
plt.grid(True)
```



In [570]: *#BirthCountry columns don't seem to be relevant. I will drop them for now.*

```
df.drop(['CountryBirthFather', 'CountryBirthMother', 'CountryBirthSelf'], axis=1, inplace=True)
df_test.drop(['CountryBirthFather', 'CountryBirthMother', 'CountryBirthSelf'], axis=1, inplace=True)
```

Citizenship

In [571]: df['Citizenship'].describe()

```
Out[571]: count      145005
          unique         5
```

```

top      Native- Born in the United States
freq                                           126782
Name: Citizenship, dtype: object

```

```
In [572]: df['Citizenship'].value_counts()
```

```

Out[572]: Native- Born in the United States      126782
Foreign born- Not a citizen of U S             10829
Foreign born- U S citizen by naturalization     4614
Native- Born in Puerto Rico or U S Outlying    1469
Native- Born abroad of American Parent(s)      1311
Name: Citizenship, dtype: int64

```

```
In [573]: df['Citizenship'].value_counts(normalize=True)
```

```

Out[573]: Native- Born in the United States      0.874328
Foreign born- Not a citizen of U S             0.074680
Foreign born- U S citizen by naturalization     0.031820
Native- Born in Puerto Rico or U S Outlying    0.010131
Native- Born abroad of American Parent(s)      0.009041
Name: Citizenship, dtype: float64


```

```


In [574]: plt.figure(figsize=(10,10))
sns.countplot(x='Income>50k', hue='Citizenship', data=df[df['Income>50k']==1],palette=
plt.grid()
plt.legend(fontsize='x-small')

```

```
Out[574]: <matplotlib.legend.Legend at 0x12da05a8cc0>
```



```
In [575]: plt.figure(figsize=(12,12))
sns.barplot(x='Citizenship', y='Income>50k', data=df, dodge=True)
plt.xticks(rotation=90)
plt.grid(True)
```




```
In [576]: #I will replace the values with 1 = US Citizen, and 0 = Non-US Citizen
df['Citizenship'] = df['Citizenship'].apply(lambda x: 0 if x == 'Foreign born- Not a'
df_test['Citizenship'] = df_test['Citizenship'].apply(lambda x: 0 if x == 'Foreign born-
```

OwnBusiness

```
In [577]: df['OwnBusiness'].value_counts()
```

```
Out[577]: 0    126718  
         2     15685  
         1       2602  
         Name: OwnBusiness, dtype: int64
```

```
In [578]: plt.figure(figsize=(12,12))  
         sns.barplot(x='OwnBusiness', y='Income>50k', data=df, dodge=True)  
         plt.xticks(rotation=90)  
         plt.grid(True)
```



```
In [579]: #Very clear correlation between Owning a Business and Income.  
         #I will rename '2' to 1.  
         df['OwnBusiness'] = np.where(df['OwnBusiness'] == 0, 0, 1)
```

```

df_test['OwnBusiness'] = np.where(df_test['OwnBusiness'] == 0, 0, 1)

df['OwnBusiness'].value_counts(normalize=True)

Out[579]: 0    0.873887
          1    0.126113
          Name: OwnBusiness, dtype: float64

VeteranBenefits'

In [580]: df['VeteranBenefits'].describe()


Out[580]: count    145005.000000
          mean       1.826165
          std        0.551837
          min        0.000000
          25%        2.000000
          50%        2.000000
          75%        2.000000
          max        2.000000
          Name: VeteranBenefits, dtype: float64

In [581]: df['VeteranBenefits'].value_counts()

Out[581]: 2    131464
          0    11666
          1    1875
          Name: VeteranBenefits, dtype: int64

In [582]: plt.figure(figsize=(12,12))
          sns.barplot(x='VeteranBenefits', y='Income>50k', data=df, dodge=True)
          plt.xticks(rotation=90)
          plt.grid(True)

```

```
In [583]: #Couldn't find any explanations for this column, what it means, what are the digits.
df.drop('VeteranBenefits', axis=1, inplace=True)
df_test.drop('VeteranBenefits', axis=1, inplace=True)
```

WeeksWorkedInY


```
In [584]: df['WeeksWorkedInY'].describe()
```

```
Out[584]: count      145005.000000
mean         30.532313
std          23.667317
min           0.000000
25%           0.000000
50%          47.000000
75%          52.000000
```

```
max          52.000000
Name: WeeksWorkedInY, dtype: float64
```


```
In [585]: plt.figure(figsize=(12,12))
          sns.boxplot(x='Income>50k', y='WeeksWorkedInY', data=df)
```

```
Out[585]: <matplotlib.axes._subplots.AxesSubplot at 0x12cc6d242e8>
```




```
In [586]: sns.distplot(df['WeeksWorkedInY'])
```

```
Out[586]: <matplotlib.axes._subplots.AxesSubplot at 0x12caab58240>
```




```
In [587]: g = sns.FacetGrid(data=df, col='Income>50k', height=6)
g.map(sns.distplot, 'WeeksWorkedInY' )
for ax in g.axes.flat:
    labels = ax.get_xticklabels() # get x labels
    ax.set_xticklabels(labels, rotation=90) # set new labels
```




```
In [588]: plt.figure(figsize=(12,12))
          sns.violinplot(x='Income>50k', y='WeeksWorkedInY', data=df)
```

```
Out [588]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4fb583c8>
```




Year

```
In [589]: plt.figure(figsize=(12,12))
          sns.barplot(x='Year', y='Income>50k', data=df, dodge=True)
          plt.xticks(rotation=90)
          plt.grid(True)
```



```
In [590]: plt.figure(figsize=(12,12))
sns.countplot(x='Year', data=df, hue='Income>50k',dodge=True)
plt.xticks(rotation=90)
plt.grid(True)
```



```
In [591]: #Year doesn't seem to be correlated with Income levels. Even the above charts were v
df.drop('Year', axis=1, inplace=True)
df_test.drop('Year', axis=1, inplace=True)
```

```
In [592]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 145005 entries, 0 to 199520
Data columns (total 20 columns):
Age                145005 non-null int64
ClassOfWorker      145005 non-null object
Education          145005 non-null object
WagePerHour        145005 non-null int64
MaritalStatus      145005 non-null object
MajorIndustryCode  145005 non-null object
```

```

MajorOccupationCode      145005 non-null object
Race                     145005 non-null object
Male                     145005 non-null int32
CapitalGains              145005 non-null int64
CapitalLosses             145005 non-null int64
StockDividends            145005 non-null int64
TaxFilerStat             145005 non-null object
HouseholdFamilyStatus     145005 non-null object
HouseholdSummary          145005 non-null object
NumPersonsWorkedEmployer  145005 non-null int64
Citizenship               145005 non-null int64
OwnBusiness               145005 non-null int32
WeeksWorkedInY           145005 non-null int64
Income>50k                145005 non-null int32
dtypes: int32(3), int64(8), object(9)
memory usage: 26.6+ MB

```

```
In [593]: df.describe()
```

```

Out [593]:

```

	Age	WagePerHour	Male	CapitalGains \
count	145005.000000	145005.000000	145005.000000	145005.000000
mean	39.418468	70.766222	0.485866	557.842136
std	19.327772	277.214096	0.499802	5286.552736
min	0.000000	0.000000	0.000000	0.000000
25%	25.000000	0.000000	0.000000	0.000000
50%	38.000000	0.000000	0.000000	0.000000
75%	52.000000	0.000000	1.000000	0.000000
max	90.000000	2000.000000	1.000000	99999.000000

	CapitalLosses	StockDividends	NumPersonsWorkedEmployer	Citizenship \
count	145005.000000	145005.000000	145005.000000	145005.000000
mean	48.684659	252.061688	2.581587	0.925320
std	308.852831	2260.621084	2.402695	0.262876
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	1.000000
50%	0.000000	0.000000	2.000000	1.000000
75%	0.000000	0.000000	5.000000	1.000000
max	4608.000000	99999.000000	6.000000	1.000000

	OwnBusiness	WeeksWorkedInY	Income>50k
count	145005.000000	145005.000000	145005.000000
mean	0.126113	30.532313	0.080577
std	0.331978	23.667317	0.272185
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	47.000000	0.000000
75%	0.000000	52.000000	0.000000
max	1.000000	52.000000	1.000000

```
In [594]: df.head()
```

```
Out[594]:
```

	Age	ClassOfWorker	Education	WagePerHour	\
0	73	NA	High school graduate	0	
1	58	Self-employed-not incorporated	High school graduate	0	
2	18	NA	Other	0	
3	9	NA	Other	0	
4	10	NA	Other	0	

	MaritalStatus	MajorIndustryCode	\
0	Widowed	Not in universe or children	
1	Divorced	Construction	
2	Never married	Not in universe or children	
3	Never married	Not in universe or children	
4	Never married	Not in universe or children	

	MajorOccupationCode	Race	Male	\
0	NA	White	0	
1	Precision production craft & repair	White	1	
2	NA	Asian or Pacific Islander	0	
3	NA	White	0	
4	NA	White	0	

	CapitalGains	CapitalLosses	StockDividends	TaxFilerStat	\
0	0	0	0	Nonfiler	
1	0	0	0	Other	
2	0	0	0	Nonfiler	
3	0	0	0	Nonfiler	
4	0	0	0	Nonfiler	

	HouseholdFamilyStatus	HouseholdSummary	\
0	Other	Other relative of householder	
1	Householder	Householder	
2	Other	Child 18 or older	
3	Other	Child under 18 never married	
4	Other	Child under 18 never married	


	NumPersonsWorkedEmployer	Citizenship	OwnBusiness	WeeksWorkedInY	\
0	0	1	0	0	
1	1	1	0	52	
2	0	0	0	0	
3	0	1	0	0	
4	0	1	0	0	

	Income>50k
0	0
1	0
2	0


```
3      0
4      0
```

```
In [595]: plt.subplots(figsize=(10, 10))
df_cor = df.corr()
sns.heatmap(df_cor, annot=True, fmt = ".1f", cmap = "coolwarm")
```

```
Out[595]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4fb709b0>
```



```
In [596]: df_test.shape
```

```
Out[596]: (74777, 20)
```

```
In [597]: df.shape
```

```
Out[597]: (145005, 20)
```

6 MODELING

6.1 Encoding

```
In [598]: df_train_encoded=pd.get_dummies(df, columns=['ClassOfWork', 'Education', 'MaritalS
          'MajorOccupationCode', 'Race', 'TaxFilerStat
          'HouseholdSummary'],
          prefix=['ClassOfWork', 'Education', 'MaritalStatus', 'Maj
          'MajorOccupationCode', 'Race', 'TaxFilerStat'
          'HouseholdSummary'])

df_test_encoded=pd.get_dummies(df_test, columns=['ClassOfWork', 'Education', 'Mari
          'MajorOccupationCode', 'Race', 'TaxFilerStat'
          'HouseholdSummary'],
          prefix=['ClassOfWork', 'Education', 'MaritalStatus', 'Maj
          'MajorOccupationCode', 'Race', 'TaxFilerStat'
          'HouseholdSummary'])

In [599]: df_train_encoded.shape

Out[599]: (145005, 86)

In [600]: X_train = df_train_encoded.loc[:,df_train_encoded.columns != 'Income>50k']
          y_train = df_train_encoded['Income>50k']
          X_test = df_test_encoded.loc[:,df_test_encoded.columns != 'Income>50k']
          y_test = df_test_encoded['Income>50k']
```

6.2 Use RandomForest for Feature Selection

```
In [601]: #First, scale the data
          from sklearn.preprocessing import StandardScaler
          scaler = StandardScaler()
          scaler.fit(X_train)
          X_train = pd.DataFrame(scaler.transform(X_train), columns=X_train.columns)
          X_test = pd.DataFrame(scaler.transform(X_test), columns=X_test.columns)

C:\Users\ahmed\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:625: DataConversionWarning: Data
  return self.partial_fit(X, y)
C:\Users\ahmed\Anaconda3\lib\site-packages\ipykernel_launcher.py:5: DataConversionWarning: Data
  """
C:\Users\ahmed\Anaconda3\lib\site-packages\ipykernel_launcher.py:6: DataConversionWarning: Data

In [602]: from sklearn.ensemble import RandomForestClassifier
          rfc = RandomForestClassifier(n_estimators=300)
          rfc.fit(X_train, y_train)

Out[602]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
          max_depth=None, max_features='auto', max_leaf_nodes=None,
```

```

min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, n_estimators=300, n_jobs=None,
oob_score=False, random_state=None, verbose=0,
warm_start=False)

```

```

In [603]: from sklearn.feature_selection import SelectFromModel
sel = SelectFromModel(rfc, prefit=True)
selected_feat= X_train.columns[(sel.get_support())]

```

```

In [604]: print('Selected Features:\n',*selected_feat, sep='\n')

```

Selected Features:

```

Age
WagePerHour
Male
CapitalGains
CapitalLosses
StockDividends
NumPersonsWorkedEmployer
OwnBusiness
WeeksWorkedInY
Education_Bachelors degree(BA AB BS)
Education_High school graduate
Education_Masters degree(MA MS MEng MEd MSW MBA)
Education_Other
Education_Prof school degree (MD DDS DVM LLB JD)
MajorOccupationCode_Executive admin and managerial
MajorOccupationCode_Professional specialty

```

```

In [605]: X_train = X_train[selected_feat]
X_test = X_test[selected_feat]

```

```

In [606]: df2 = pd.concat([X_train,y_train.reset_index()], axis=1)
df2.drop('index', axis = 1 ,inplace=True)

```

```


plt.subplots(figsize=(10, 10))
df_cor = df2.corr()
sns.heatmap(df_cor, annot=True, fmt = ".1f", cmap = "coolwarm")

```

```

Out[606]: <matplotlib.axes._subplots.AxesSubplot at 0x12d4eb89898>

```



6.3 Model Selection

In [607]: `X_train.shape`

Out[607]: (145005, 16)

```
In [608]: from sklearn.tree import DecisionTreeClassifier
           from sklearn.linear_model import LogisticRegression
           from sklearn.svm import SVC
           from sklearn.model_selection import GridSearchCV, cross_val_score, StratifiedKFold, ...
```

In [609]: `kfold = StratifiedKFold(n_splits=10)`

```
In [610]: #Doing 10-fold cross validation, using Decision Tree and Logistic Regression
```

```
rs=42
```

```
classifiers = [] # list of classifiers tested
```

```
classifiers.append(LogisticRegression(random_state = rs))
```

```
classifiers.append(DecisionTreeClassifier(random_state = rs))
```

```
cv_results = []
```

```
for classifier in classifiers :
```

```
    cv_results.append(cross_val_score(classifier, X_train, y_train, scoring = 'roc_auc'))
```

```
cv_means = []
```

```
cv_std = []
```

```
for cv_result in cv_results:
```

```
    cv_means.append(cv_result.mean())
```

```
    cv_std.append(cv_result.std())
```

```
cv_res = pd.DataFrame({'CV_score':cv_means, 'CV_stddev':cv_std, 'Algorithm':['LogisticRegression', 'DecisionTree']})
```

```
In [611]: cv_res
```

```
Out[611]:
```


	CV_score	CV_stddev	Algorithm
0	0.908088	0.006893	LogisticRegression
1	0.729347	0.005867	DecisionTree

```
In [612]: plt.subplots(figsize=(10, 10))
```

```
g = sns.barplot('CV_score', 'Algorithm', data = cv_res, palette='Set2', orient = 'h',
```

```
g.set_xlabel('Mean AUC score')
```

```
g = g.set_title('Cross validation scores')
```



6.4 Logistic Regression

```
In [613]: from sklearn.metrics import roc_curve, auc
```


```
LR = LogisticRegression(random_state=42)
y_score = LR.fit(X_train, y_train).decision_function(X_test)
fpr = dict()
tpr = dict()
roc_auc = dict()

fpr[0], tpr[0], _ = roc_curve(y_test, y_score)
roc_auc[0] = auc(fpr[0], tpr[0])

# Compute micro-average ROC curve and ROC area
fpr[1], tpr[1], _ = roc_curve(y_test.ravel(), y_score.ravel())
roc_auc[1] = auc(fpr[1], tpr[1])
```

C:\Users\ahmed\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning
FutureWarning)

```
In [614]: plt.figure(figsize=(10,10))  
         lw = 2  
         plt.plot(fpr[0], tpr[0], color='darkorange',  
                  lw=lw, label='ROC curve (area = %0.2f)' % roc_auc[0])  
         plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')  
         plt.xlim([0.0, 1.0])  
         plt.ylim([0.0, 1.05])  
         plt.xlabel('False Positive Rate')  
         plt.ylabel('True Positive Rate')  
         plt.title('ROC - Logistic Regression')  
         plt.legend(loc="lower right")  
         plt.show()
```



```

In [615]: import itertools
          from sklearn.metrics import confusion_matrix

          def plot_confusion_matrix(cm, classes,
                                   normalize=False,
                                   title='Confusion matrix',
                                   cmap=plt.cm.Blues):
              """
              This function prints and plots the confusion matrix.
              Normalization can be applied by setting `normalize=True`.
              """
              if normalize:
                  cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
                  print("Normalized confusion matrix")
              else:
                  print('Confusion matrix, without normalization')

              print(cm)

              plt.imshow(cm, interpolation='nearest', cmap=cmap)
              plt.title(title)
              plt.colorbar()
              tick_marks = np.arange(len(classes))
              plt.xticks(tick_marks, classes, rotation=45)
              plt.yticks(tick_marks, classes)

              fmt = '.2f' if normalize else 'd'
              thresh = cm.max() / 2.
              for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
                  plt.text(j, i, format(cm[i, j], fmt),
                           horizontalalignment="center",
                           color="white" if cm[i, j] > thresh else "black")

              plt.tight_layout()
              plt.ylabel('True label')
              plt.xlabel('Predicted label')

In [616]: # use trained model to make predictions on test set
          y_pred = LR.predict(X_test)

In [617]: # Compute confusion matrix
          cnf_matrix = confusion_matrix(y_test, y_pred)
          np.set_printoptions(precision=2)

          class_names = ['Income < 50k', 'Income > 50k']

```



```
# Plot non-normalized confusion matrix
plt.figure(figsize=(10,10))
plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix (counts)')


# Plot normalized confusion matrix
plt.figure(figsize=(10,10))
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Confusion matrix (percent)')
```


Confusion matrix, without normalization

```
[[68137  813]
 [ 3901 1926]]
```

Normalized confusion matrix

```
[[ 0.99  0.01]
 [ 0.67  0.33]]
```





```
In [618]: from sklearn.metrics import accuracy_score, precision_score, recall_score

print("Accuracy: " , "%.2f" % (accuracy_score(y_test, y_pred)*100), '%')
print("Precision: " , "%.2f" % (precision_score(y_test, y_pred)*100), '%')
print("Recall: " , "%.2f" % (recall_score(y_test, y_pred)*100), '%')
```

Accuracy: 93.70 %
Precision: 70.32 %
Recall: 33.05 %

Accuracy is 93.70%, i.e. better than simply predicting that everybody make <50k.
Precision 70.32% means that 70.32 of the one predicted >50k are actually >50k

Recall 33% means that out of the ones who actually make >50k, the model could only find 33% of them.

Given the skewness of the data, I will to use Oversampling, hoping to get better results.

6.5 Oversampling

Due to the unbalanced data I use SMOTE to oversample the training data.

```
In [619]: from imblearn.over_sampling import SMOTE
```

```
print("Number transactions X_train dataset: ", X_train.shape)
print("Number transactions y_train dataset: ", y_train.shape)
print("Number transactions X_test dataset: ", X_test.shape)
print("Number transactions y_test dataset: ", y_test.shape)
```

```
Number transactions X_train dataset: (145005, 16)
Number transactions y_train dataset: (145005,)
Number transactions X_test dataset: (74777, 16)
Number transactions y_test dataset: (74777,)
```

```
In [620]: print("Before OverSampling, counts of label '1': {}".format(sum(y_train==1)))
print("Before OverSampling, counts of label '0': {} \n".format(sum(y_train==0)))

sm = SMOTE(random_state=2)
X_train_res, y_train_res = sm.fit_sample(X_train, y_train.ravel())

print('After OverSampling, the shape of train_X: {}'.format(X_train_res.shape))
print('After OverSampling, the shape of train_y: {} \n'.format(y_train_res.shape))

print("After OverSampling, counts of label '1': {}".format(sum(y_train_res==1)))
print("After OverSampling, counts of label '0': {}".format(sum(y_train_res==0)))
```

```
Before OverSampling, counts of label '1': 11684
Before OverSampling, counts of label '0': 133321
```

```
After OverSampling, the shape of train_X: (266642, 16)
After OverSampling, the shape of train_y: (266642,)
```

```
After OverSampling, counts of label '1': 133321
After OverSampling, counts of label '0': 133321
```

```
In [621]: X_train_res = pd.DataFrame(X_train_res, columns = X_train.columns)
```

```
In [622]: #Using Random Forest for feature selection
rfc.fit(X_train_res, y_train_res)
sel = SelectFromModel(rfc, prefit=True)
selected_feat= X_train_res.columns[(sel.get_support())]
```

```

print('Selected Features:\n',*selected_feat, sep='\n')
X_train_res = X_train_res[selected_feat]
X_test = X_test[selected_feat]

```

Selected Features:

```

Age
Male
StockDividends
NumPersonsWorkedEmployer
WeeksWorkedInY

```

In [623]: *#Doing 10-fold cross validation, using Decision Tree and Logistic Regression*

```

rs=42

classifiers = [] # list of classifiers tested
classifiers.append(LogisticRegression(random_state = rs))
classifiers.append(DecisionTreeClassifier(random_state = rs))

cv_results = []
for classifier in classifiers :
    cv_results.append(cross_val_score(classifier, X_train_res, y_train_res, scoring = 'roc_auc'))

cv_means = []
cv_std = []
for cv_result in cv_results:
    cv_means.append(cv_result.mean())
    cv_std.append(cv_result.std())

cv_res = pd.DataFrame({'CV_score':cv_means, 'CV_stddev':cv_std, 'Algorithm':['LogisticRegression', 'DecisionTreeClassifier']})

```

In [624]: cv_res

```


Out[624]:   CV_score  CV_stddev  Algorithm
0  0.854965  0.006464  LogisticRegression
1  0.936062  0.060799  DecisionTree

```

```

In [625]: plt.subplots(figsize=(10, 10))
g = sns.barplot('CV_score', 'Algorithm', data = cv_res, palette='Set2', orient = 'h',
               )
g.set_xlabel('Mean AUC score')
g = g.set_title('Cross validation scores')

```



```
In [626]: # define LogisticRegression
```

```
LR = DecisionTreeClassifier(random_state=42)
```

```
# fit LR model to (oversampled) training data
```

```
LR.fit(X_train_res, y_train_res)
```

```
Out [626]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=42,
splitter='best')
```

```
In [627]: # use trained model to make predictions on test set
```

```
y_pred = LR.predict(X_test)
```

```
In [628]: # Compute confusion matrix
```

```
cnf_matrix = confusion_matrix(y_test, y_pred)
```

```

np.set_printoptions(precision=2)


class_names = ['Income < 50k', 'Income > 50k']


# Plot non-normalized confusion matrix
plt.figure(figsize=(10,10))
plot_confusion_matrix(cnf_matrix, classes=class_names,
                      title='Confusion matrix (counts)')

# Plot normalized confusion matrix
plt.figure(figsize=(10,10))
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True,
                      title='Confusion matrix (percent)')

Confusion matrix, without normalization
[[64303  4647]
 [ 3499  2328]]
Normalized confusion matrix
[[ 0.93  0.07]
 [ 0.6   0.4 ]]

```





```
In [629]: print("Accuracy: " , "%.2f" % (accuracy_score(y_test, y_pred)*100), '%')
          print("Precision: " , "%.2f" % (precision_score(y_test, y_pred)*100), '%')
          print("Recall: " , "%.2f" % (recall_score(y_test, y_pred)*100), '%')
```

```
Accuracy:  89.11 %
Precision:  33.38 %
Recall:    39.95 %
```

Recall rate got much better after oversampling, but accuracy and precision went down.

6.6 Conclusion

Based on the above two scenarios (with and without oversampling), the user can select the one that matches his requirements. Since Recall and precision varied significantly, then it will depend on what is more important:

- 1- Making correct predictions => No oversampling
- 2- Finding as many >50k as possible => Oversampling
- 3- Correctly predicting >50k => No Oversampling